

Transfer Learning for Optimal Configuration of Big Data Software

Pooyan Jamshidi, Giuliano Casale, Salvatore Dipietro

Imperial College London

p.jamshidi@imperial.ac.uk

Department of Computing
Research Associate Symposium
14th June 2016

**Imperial College
London**

Motivation

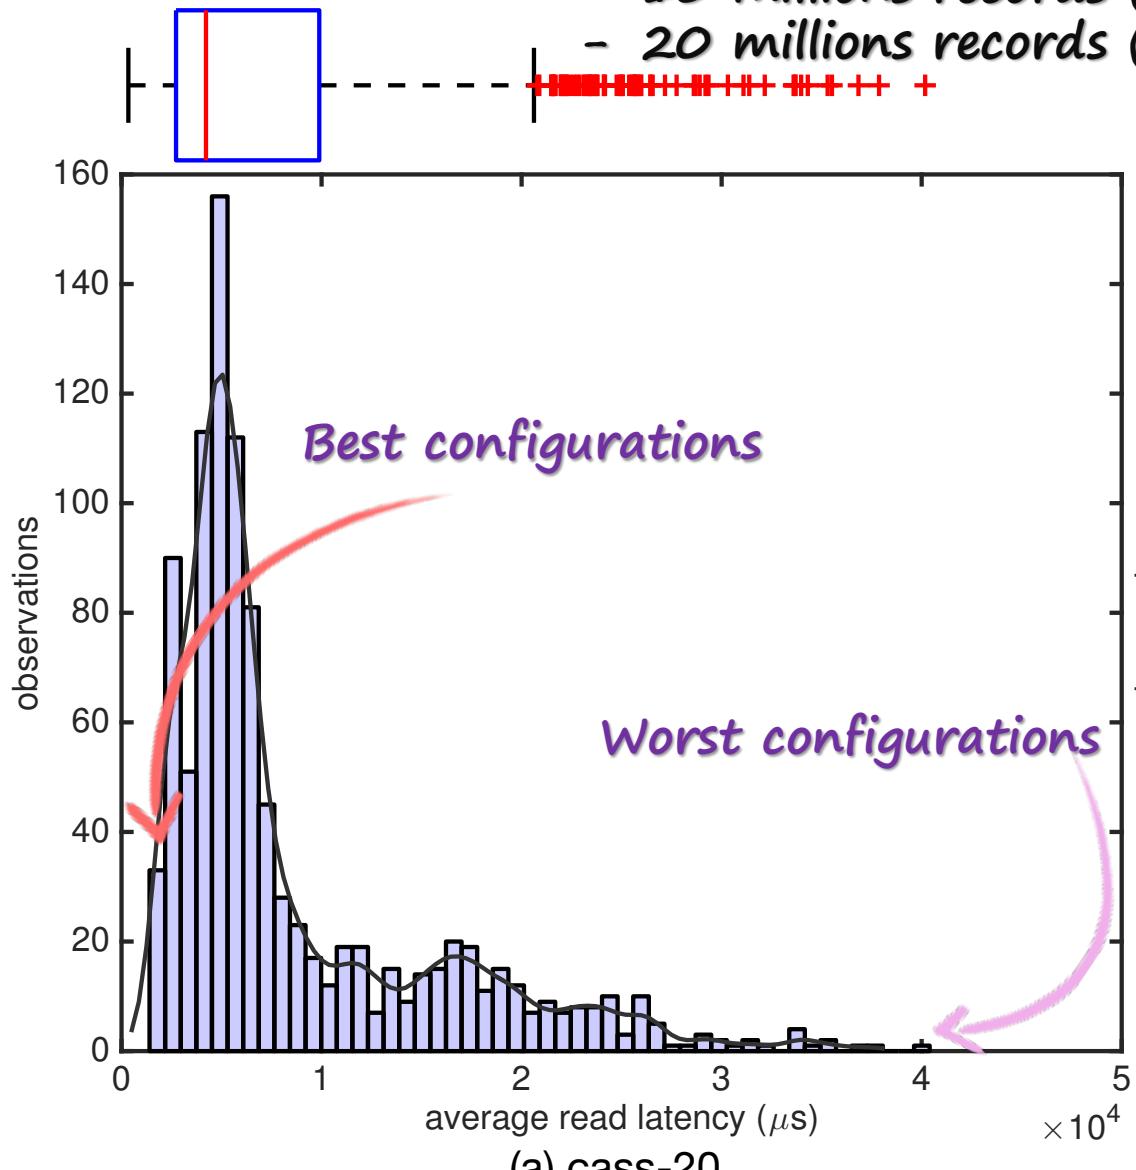
- 1- Many different Parameters =>
 - large state space
 - interactions
- 2- Defaults are typically used =>
 - poor performance

```
102 drpc.port: 3772
103 drpc.worker.threads: 64
104 drpc.max_buffer_size: 1048576
105 drpc.queue.size: 128
106 drpc.invocations.port: 3773
107 drpc.invocations.threads: 64
108 drpc.request.timeout.secs: 600
109 drpc.childopts: "-Xmx768m"
110 drpc.http.port: 3774
111 drpc.https.port: -1
112 drpc.https.keystore.password: ""
113 drpc.https.keystore.type: "JKS"
114 drpc.http.creds.plugin: org.apache.storm.security.auth.DefaultHttpCredentialsPlugin
115 drpc.authorizer.acl.filename: "drpc-auth-acl.yaml"
116 drpc.authorizer.acl.strict: false
117
118 transactional.zookeeper.root: "/transactional"
119 transactional.zookeeper.servers: null
120 transactional.zookeeper.port: null
121
122 ## blobstore configs
123 supervisor.blobstore.class: "org.apache.storm.blobstore.NimbusBlobStore"
124 supervisor.blobstore.download.thread.count: 5
125 supervisor.blobstore.download.max_retries: 3
126 supervisor.localizer.cache.target.size.mb: 10240
127 supervisor.localizer.cleanup.interval.ms: 600000
128
129
```

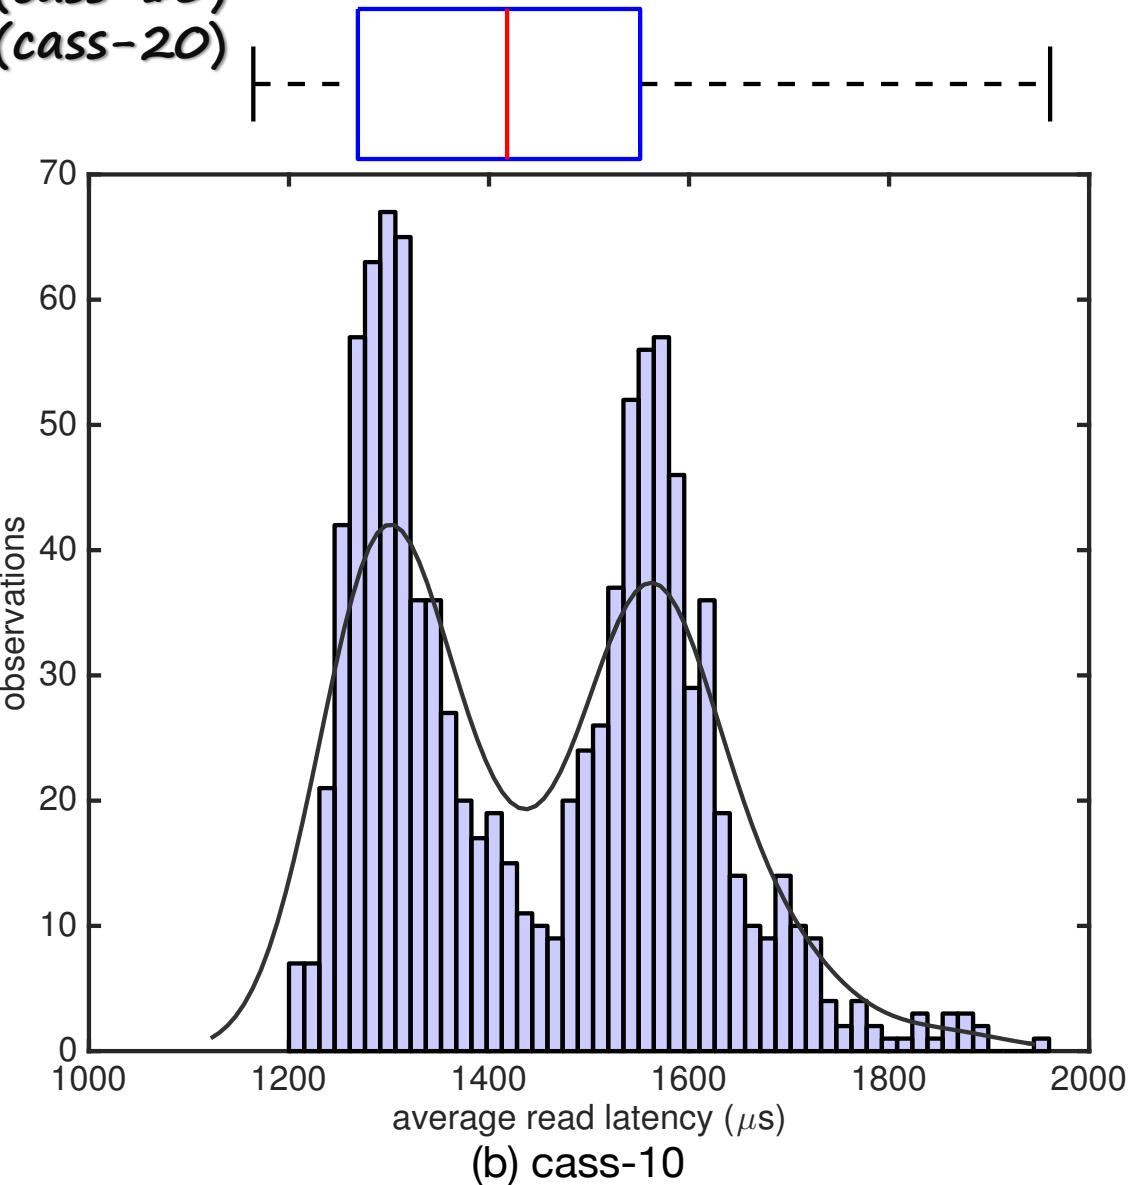
Motivation

Experiments on Apache Cassandra:

- 6 parameters, 1024 configurations
- Average read latency
- 10 millions records (cass-10)
- 20 millions records (cass-20)



(a) cass-20



(b) cass-10

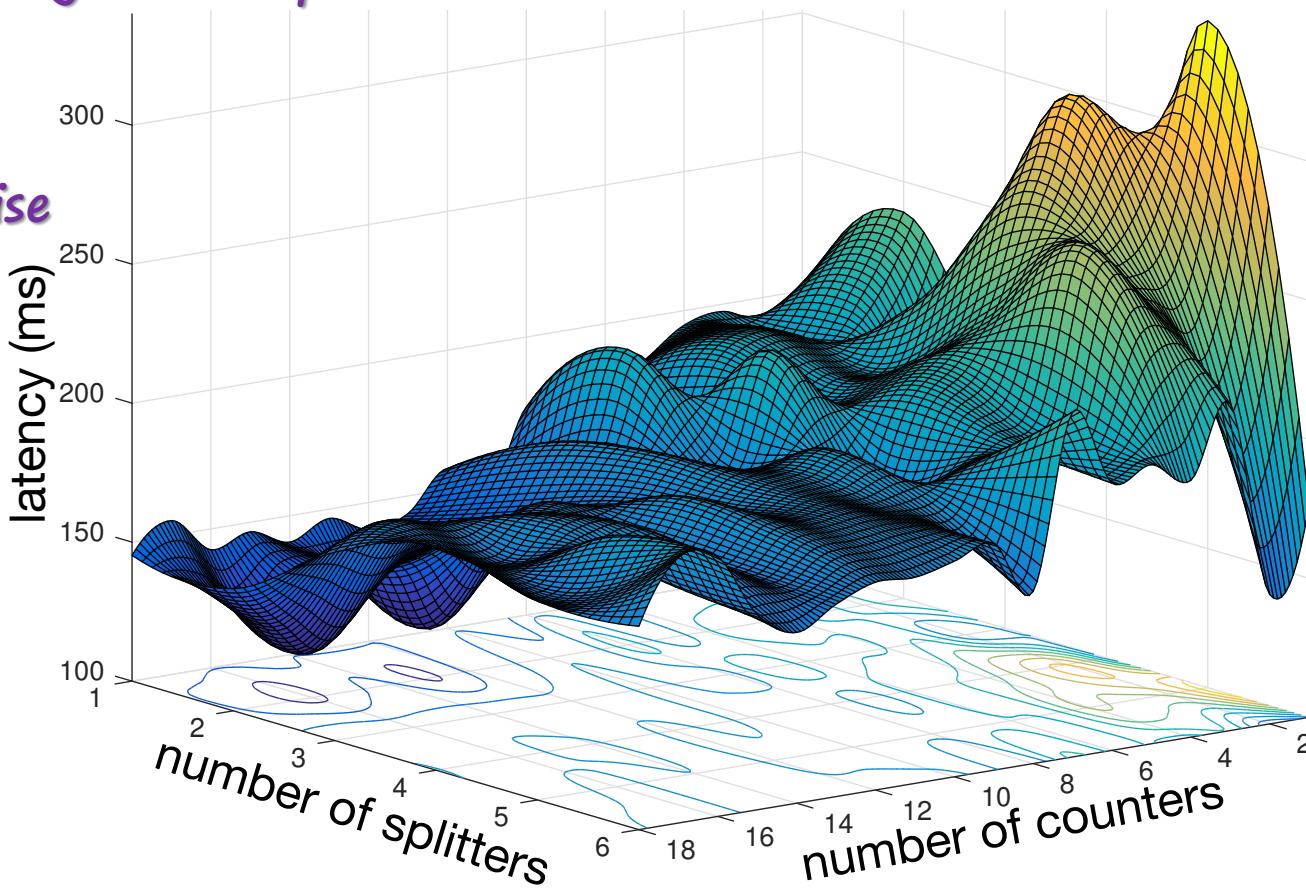
Goal!

$$\boldsymbol{x}^* = \arg \min_{\boldsymbol{x} \in \mathbb{X}} f(\boldsymbol{x})$$

$\mathbb{X} = Dom(X_1) \times \cdots \times Dom(X_d)$ Configuration space

$y_i = f(\boldsymbol{x}_i), \boldsymbol{x}_i \subset \mathbb{X}$ Partially known

$y_i = f(\boldsymbol{x}_i) + \epsilon$ Measurements contain noise



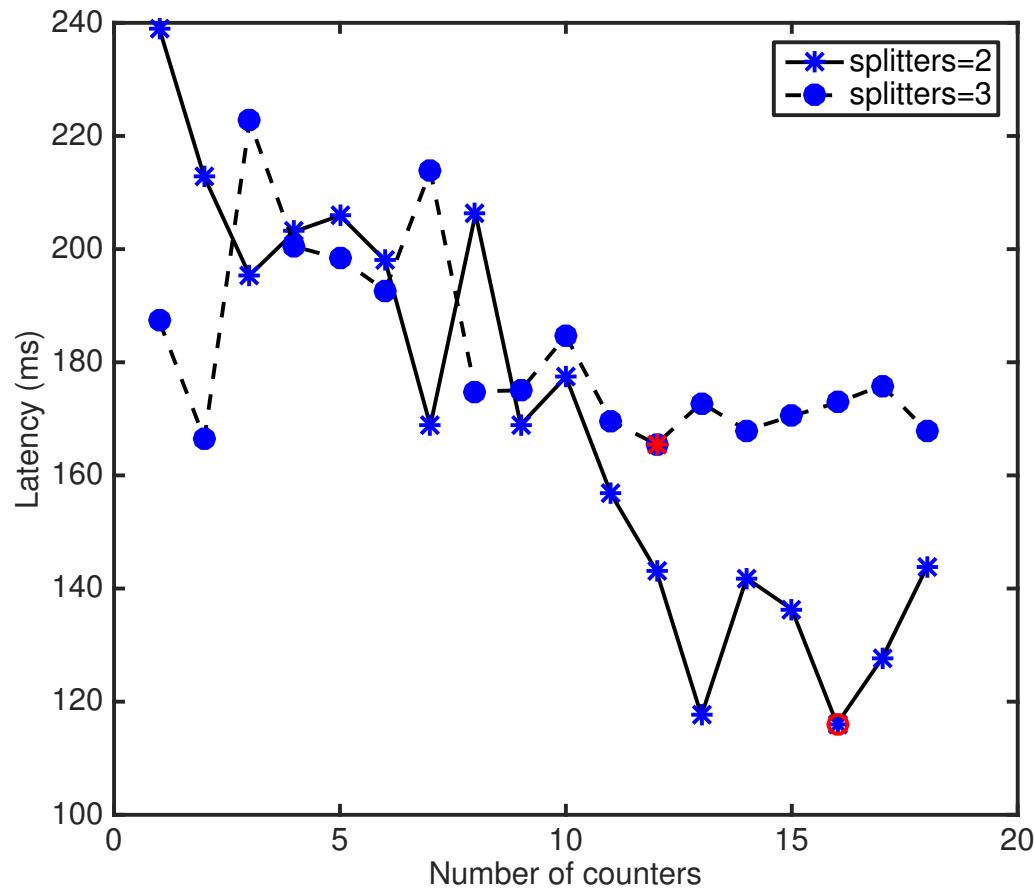
Finding optimum configuration is difficult!

$$\mathbf{x}^* = \arg \min_{\mathbf{x} \in \mathbb{X}} f(\mathbf{x})$$

$$\mathbb{X} = Dom(X_1) \times \cdots \times Dom(X_d)$$

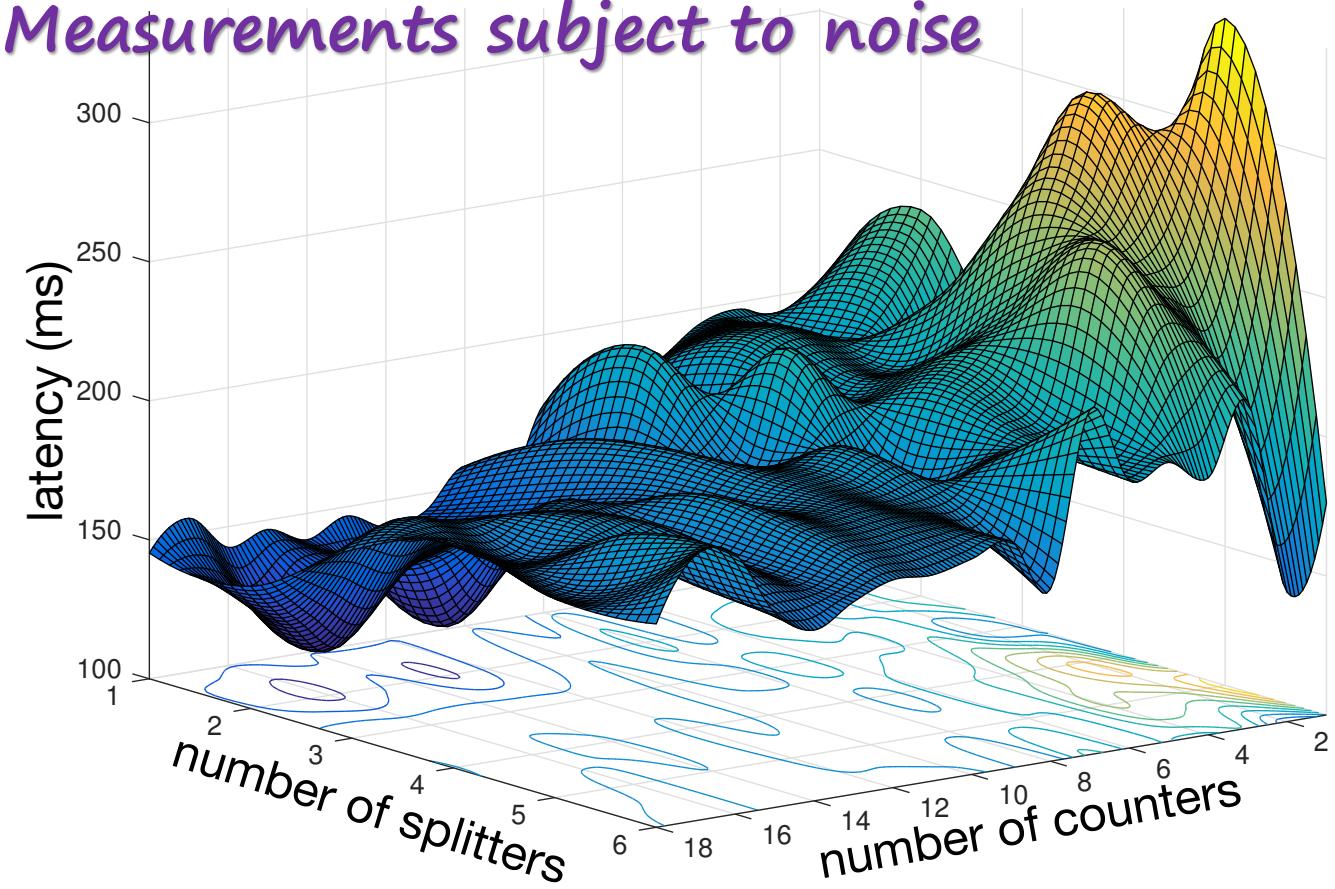
$$y_i = f(\mathbf{x}_i), \mathbf{x}_i \subset \mathbb{X}$$

$$y_i = f(\mathbf{x}_i) + \epsilon$$



Response surface is:

- Non-linear
- Non convex
- Multi-modal
- Measurements subject to noise





Part

Bayesian Optimization for Configuration Optimization (BO4CO)

Code+Data: <https://github.com/dice-project/DICE-Configuration-BO4CO>

Paper: P. Jamshidi, G. Casale, “An Uncertainty-Aware Approach to Optimal Configuration of Stream Processing Systems”, MASCOTS 2016.

<https://www.doc.ic.ac.uk/~pjamshid/PDF/mascots16.pdf>

GP for modeling blackbox response function

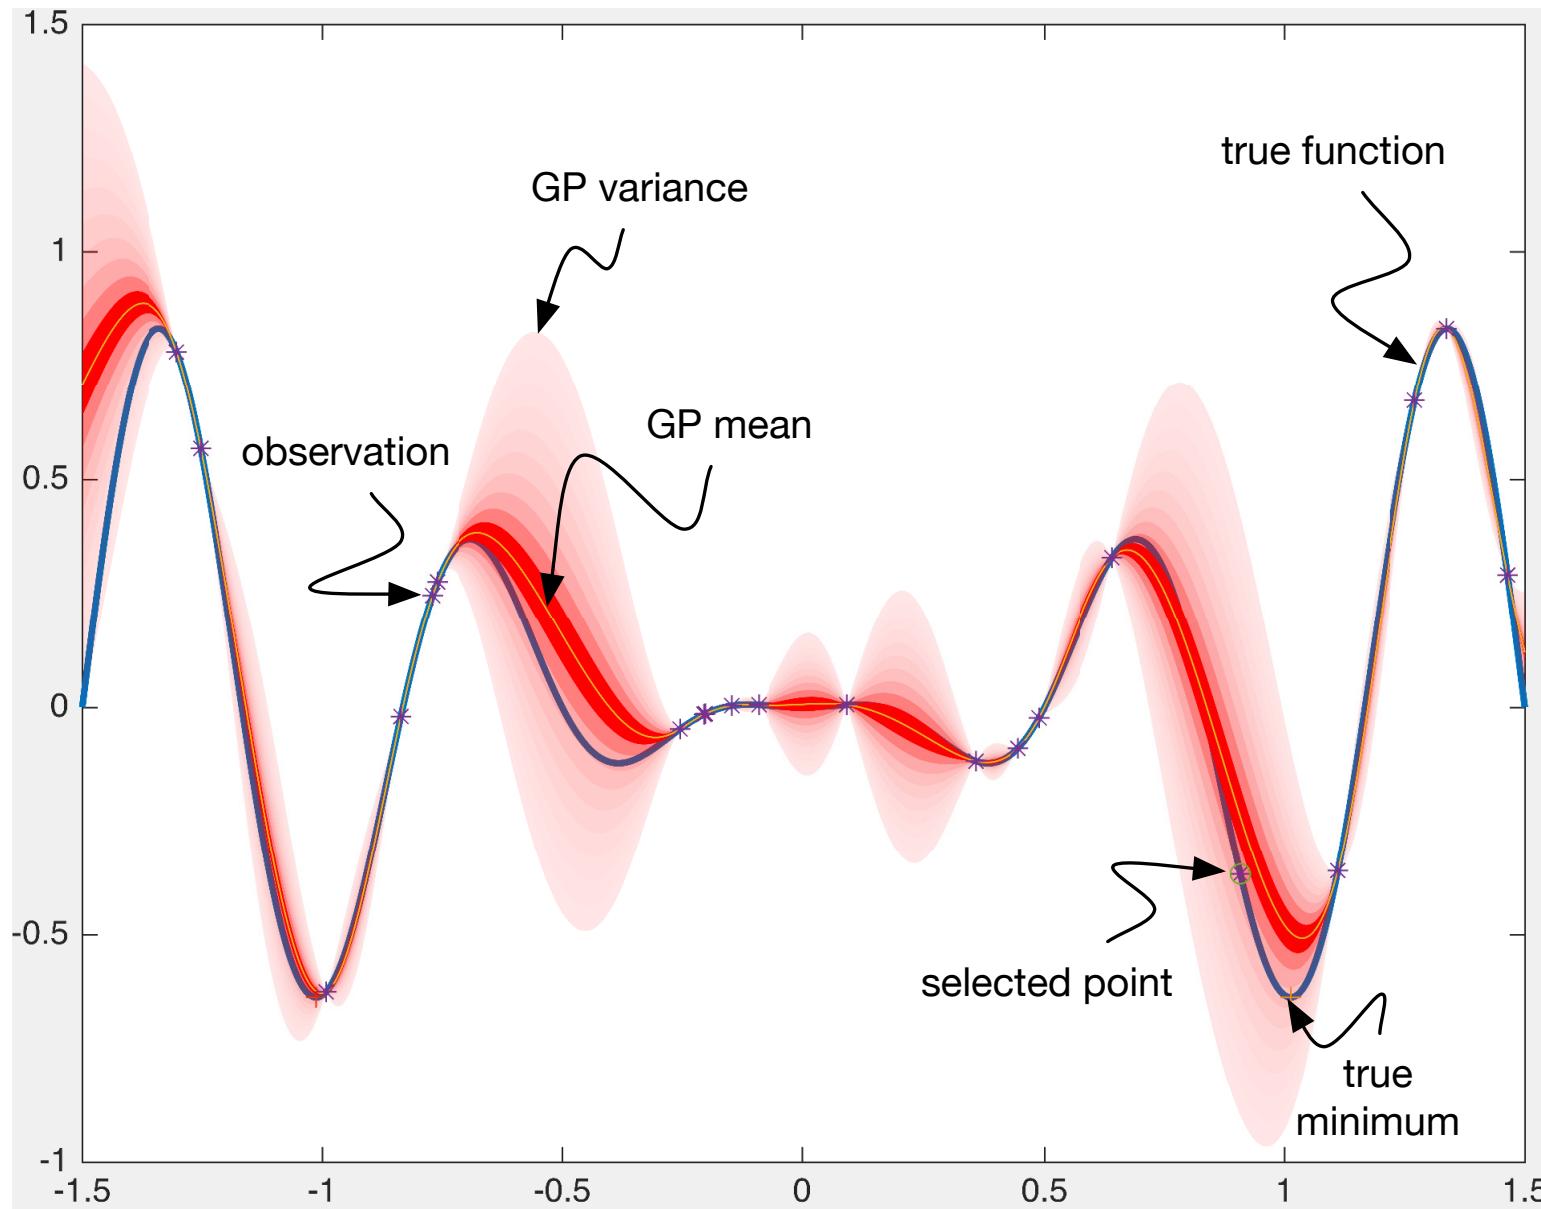
$$y = f(\mathbf{x}) \sim \mathcal{GP}(\mu(\mathbf{x}), k(\mathbf{x}, \mathbf{x}')),$$

$$\mu_t(\mathbf{x}) = \mu(\mathbf{x}) + \mathbf{k}(\mathbf{x})^\top (\mathbf{K} + \sigma^2 \mathbf{I})^{-1} (\mathbf{y} - \mu)$$

$$\sigma_t^2(\mathbf{x}) = k(\mathbf{x}, \mathbf{x}) + \sigma^2 \mathbf{I} - \mathbf{k}(\mathbf{x})^\top (\mathbf{K} + \sigma^2 \mathbf{I})^{-1} \mathbf{k}(\mathbf{x})$$

Motivations:

- 1- mean estimates + variance
- 2- all computations are linear algebra



Kernel function:

$$K := \begin{bmatrix} k(\mathbf{x}_1, \mathbf{x}_1) & \dots & k(\mathbf{x}_1, \mathbf{x}_t) \\ \vdots & \ddots & \vdots \\ k(\mathbf{x}_t, \mathbf{x}_1) & \dots & k(\mathbf{x}_t, \mathbf{x}_t) \end{bmatrix}$$

$$k_\theta(\mathbf{x}_i, \mathbf{x}_j) = \exp(\sum_{\ell=1}^d (-\theta_\ell \delta(\mathbf{x}_i \neq \mathbf{x}_j))),$$

Acquisition function:

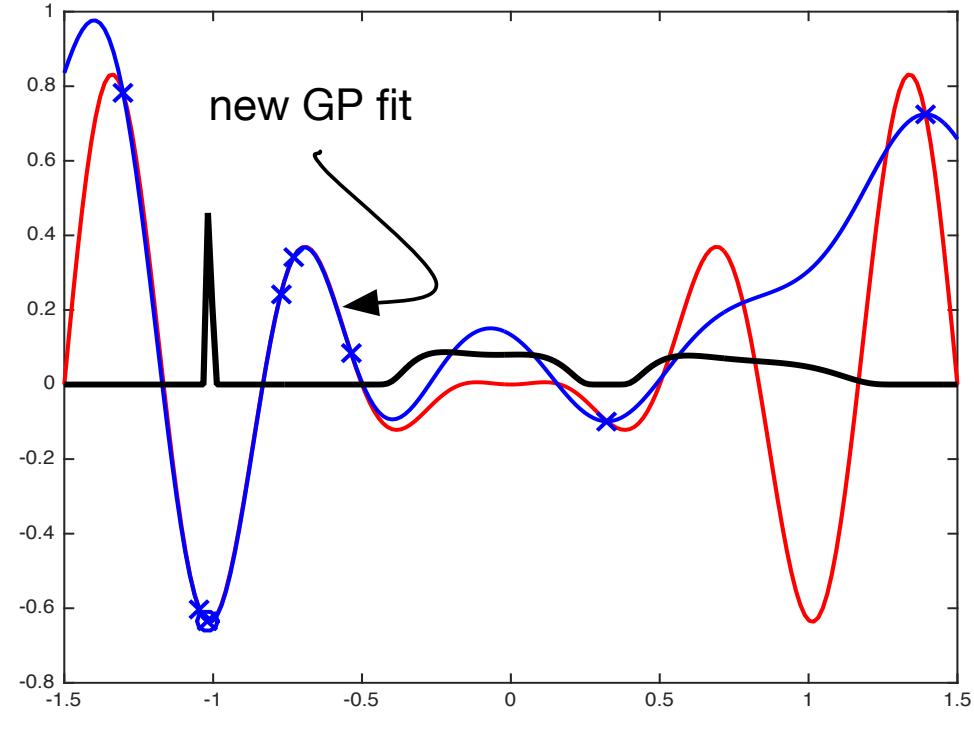
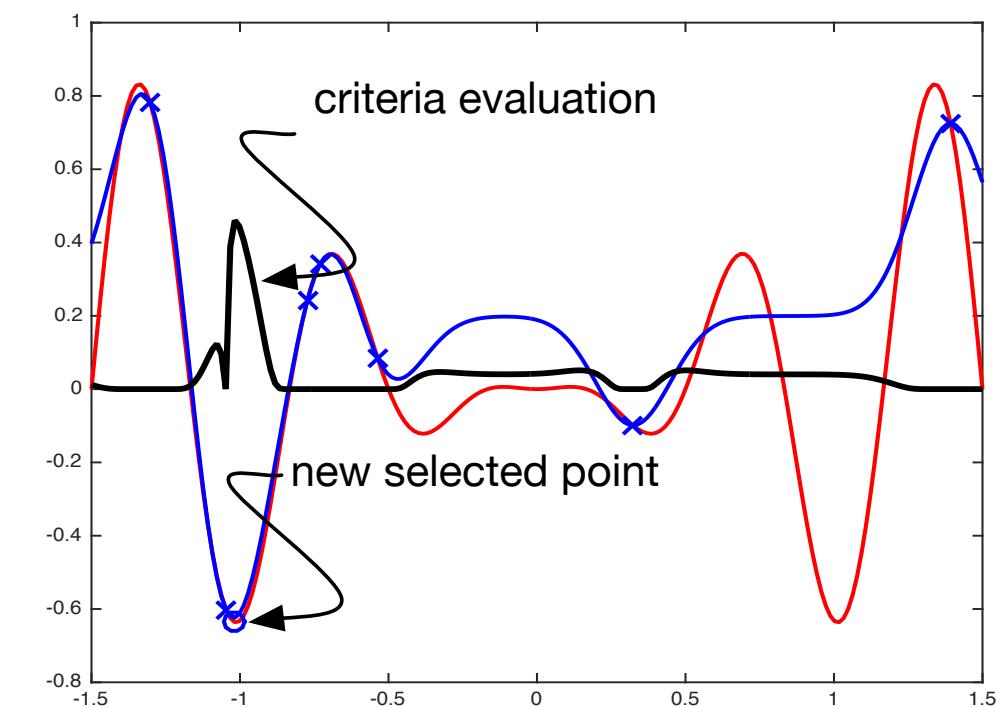
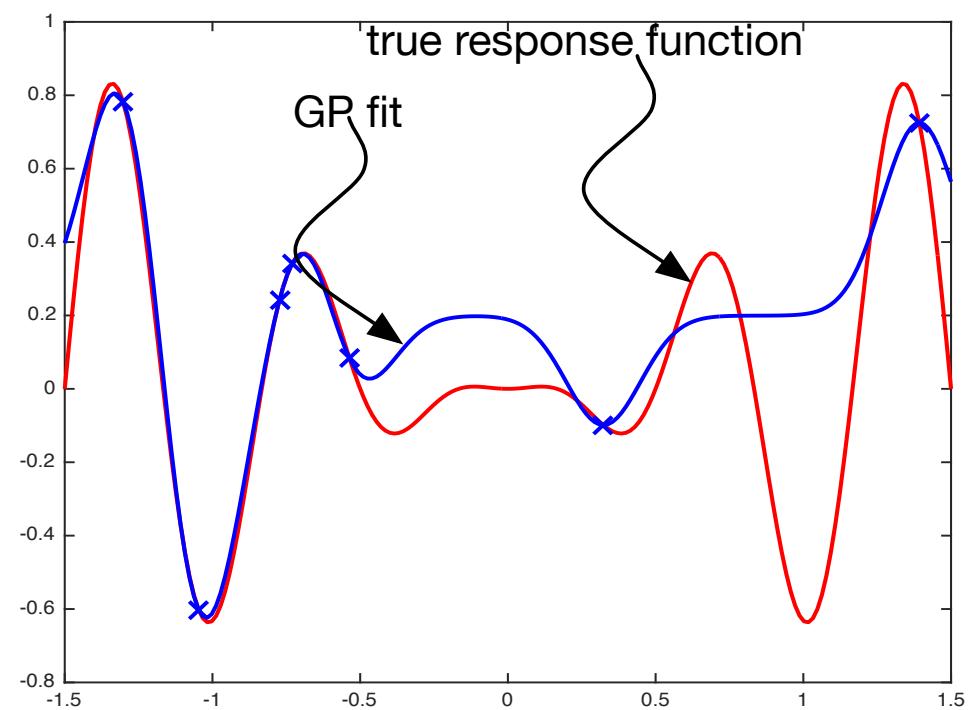
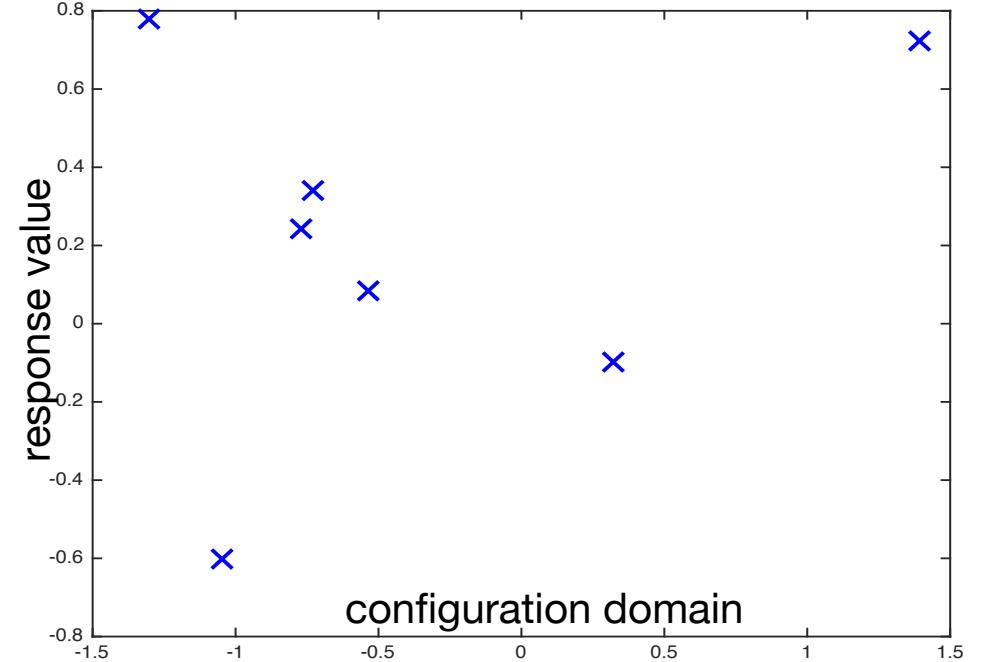
$$u_{LCB}(\mathbf{x} | \mathcal{M}, \mathbb{S}_{1:n}) = \operatorname{argmin}_{\mathbf{x} \in \mathbb{X}} \mu_t(\mathbf{x}) - \kappa \sigma_t(\mathbf{x}),$$

Algorithm 1 : BO4CO

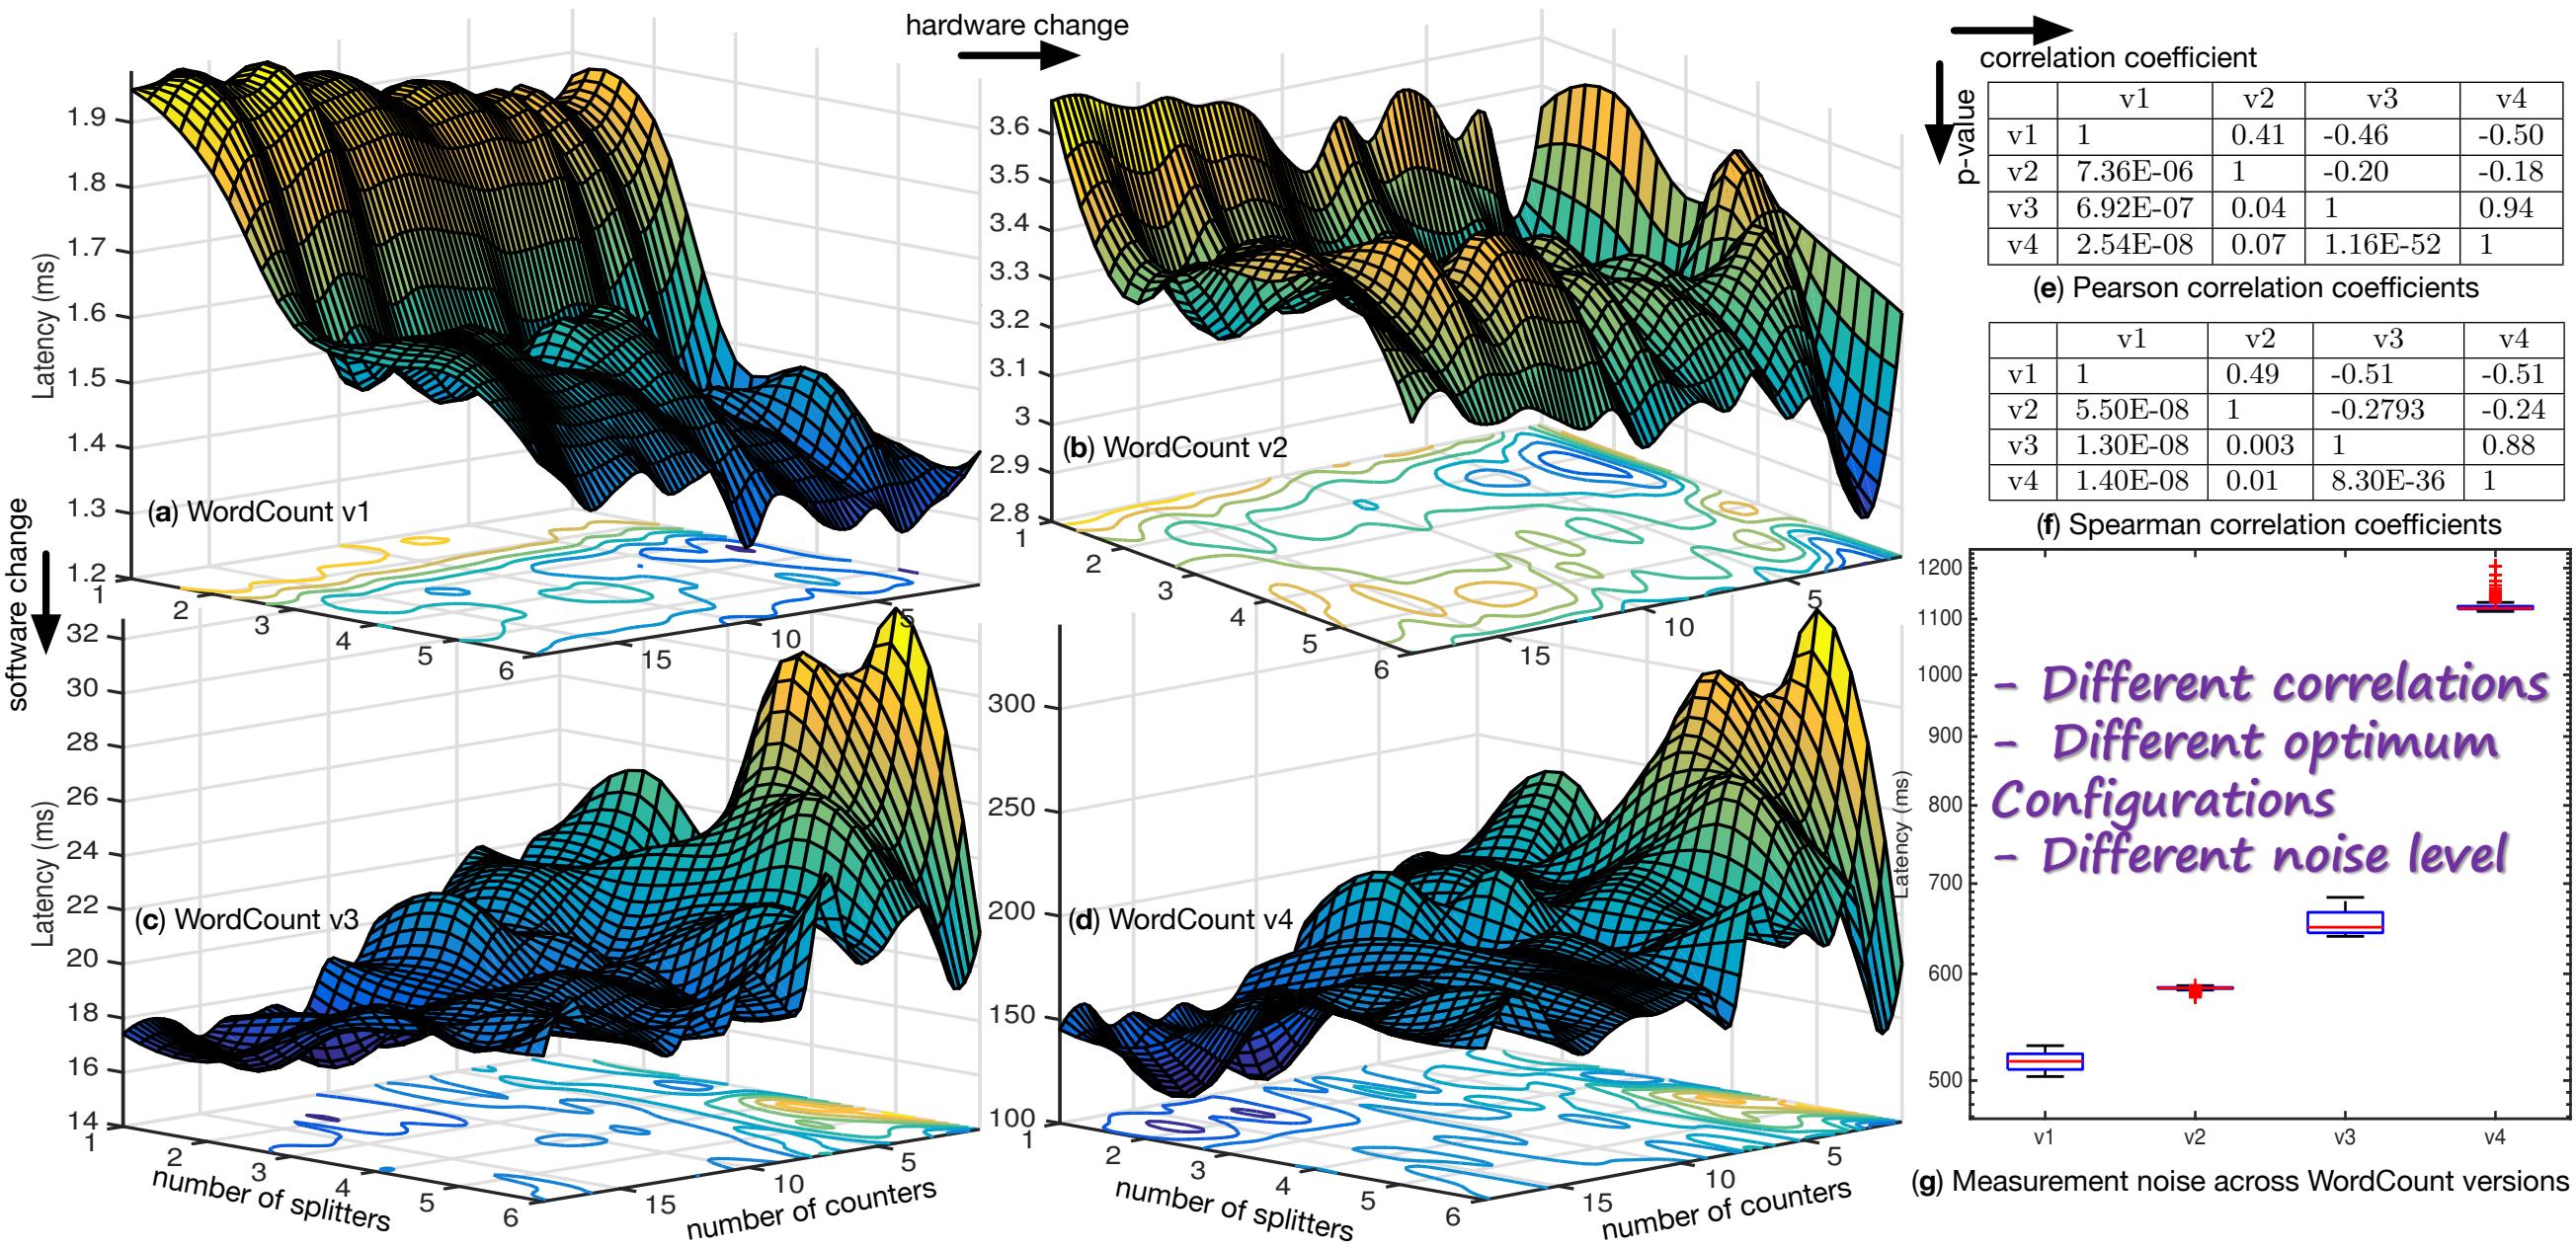
Input: Configuration space \mathbb{X} , Maximum budget N_{max} , Response function f , Kernel function K_θ , Hyper-parameters θ , Design sample size n , learning cycle N_l

Output: Optimal configurations \mathbf{x}^* and learned model \mathcal{M}

- 1: choose an initial sparse design (*lhd*) to find an initial design samples $\mathcal{D} = \{\mathbf{x}_1, \dots, \mathbf{x}_n\}$
 - 2: obtain *performance measurements* of the initial design, $y_i \leftarrow f(\mathbf{x}_i) + \epsilon_i, \forall \mathbf{x}_i \in \mathcal{D}$
 - 3: $\mathbb{S}_{1:n} \leftarrow \{(\mathbf{x}_i, y_i)\}_{i=1}^n; t \leftarrow n + 1$
 - 4: $\mathcal{M}(\mathbf{x} | \mathbb{S}_{1:n}, \theta) \leftarrow$ fit a \mathcal{GP} model to the design \triangleright Eq.(3)
 - 5: **while** $t \leq N_{max}$ **do**
 - 6: if $(t \bmod N_l = 0)$ $\theta \leftarrow$ learn the kernel hyper-parameters by maximizing the likelihood
 - 7: find *next configuration* \mathbf{x}_t by optimizing the selection criteria over the estimated response surface given the data, $\mathbf{x}_t \leftarrow \arg \max_{\mathbf{x}} u(\mathbf{x} | \mathcal{M}, \mathbb{S}_{1:t-1})$ \triangleright Eq.(9)
 - 8: obtain performance for the *new configuration* \mathbf{x}_t , $y_t \leftarrow f(\mathbf{x}_t) + \epsilon_t$
 - 9: Augment the configuration $\mathbb{S}_{1:t} = \{\mathbb{S}_{1:t-1}, (\mathbf{x}_t, y_t)\}$
 - 10: $\mathcal{M}(\mathbf{x} | \mathbb{S}_{1:t}, \theta) \leftarrow$ re-fit a new GP model \triangleright Eq.(7)
 - 11: $t \leftarrow t + 1$
 - 12: **end while**
 - 13: $(\mathbf{x}^*, y^*) = \min \mathbb{S}_{1:N_{max}}$
 - 14: $\mathcal{M}(\mathbf{x})$
-



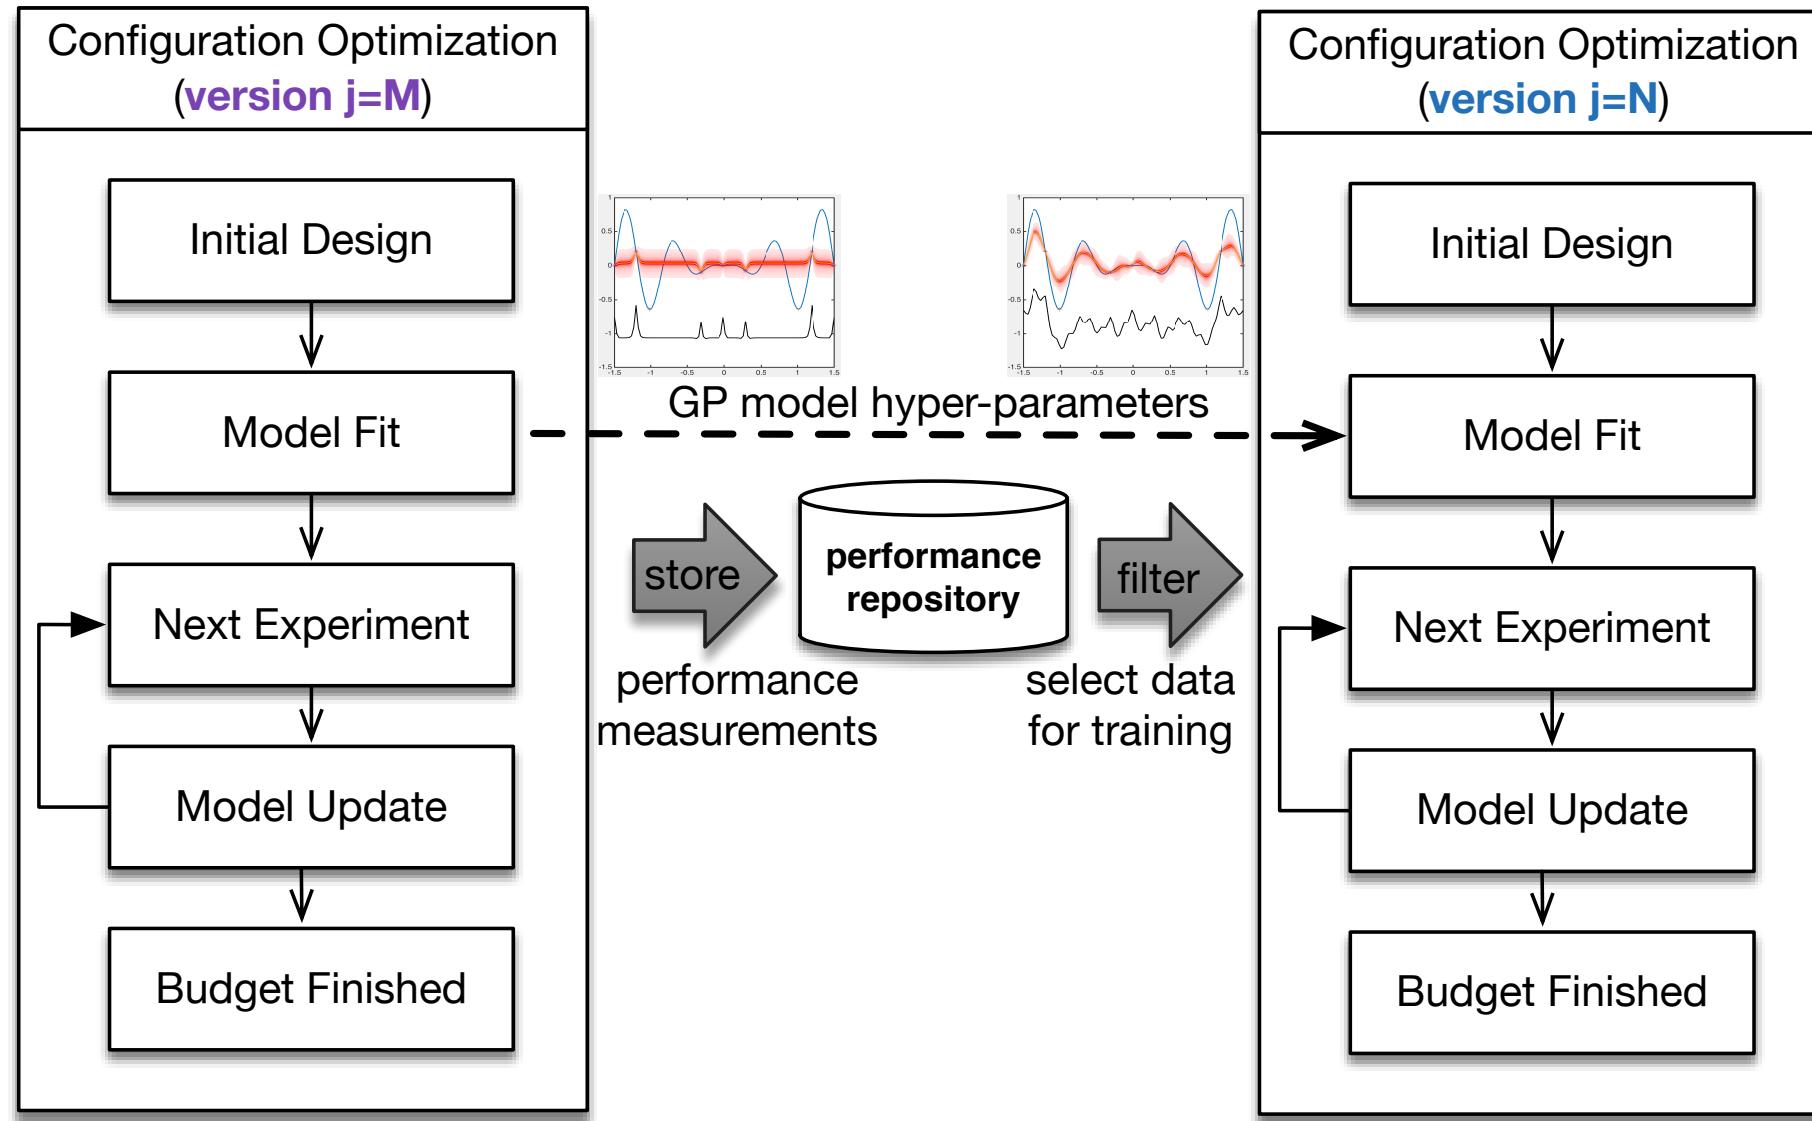
Correlations across different versions



DevOps

- Different versions are continuously delivered (daily basis).
- Big Data systems are developed using similar frameworks (Apache Storm, Spark, Hadoop, Kafka, etc).
- Different versions share similar business logics.

Solution: Transfer Learning for Configuration Optimization





Part

Transfer Learning for Configuration Optimization (TL4CO)

Code: <https://github.com/dice-project/DICE-Configuration-TL4CO>

$$\mu_t(\mathbf{x}) = \mu(\mathbf{x}) + \mathbf{k}^\top (\mathbf{K}(X, X) + \Sigma)^{-1} (\mathbf{y} - \boldsymbol{\mu})$$

$$\sigma_t^2(\mathbf{x}) = k(\mathbf{x}, \mathbf{x}) + \sigma^2 - \mathbf{k}^\top (\mathbf{K}(X, X) + \Sigma)^{-1} \mathbf{k},$$

$$k_{TL4CO}(l, l', \mathbf{x}, \mathbf{x}') = k_t(l, l') \times k_{xx}(\mathbf{x}, \mathbf{x}'),$$

correlation between different versions covariance functions

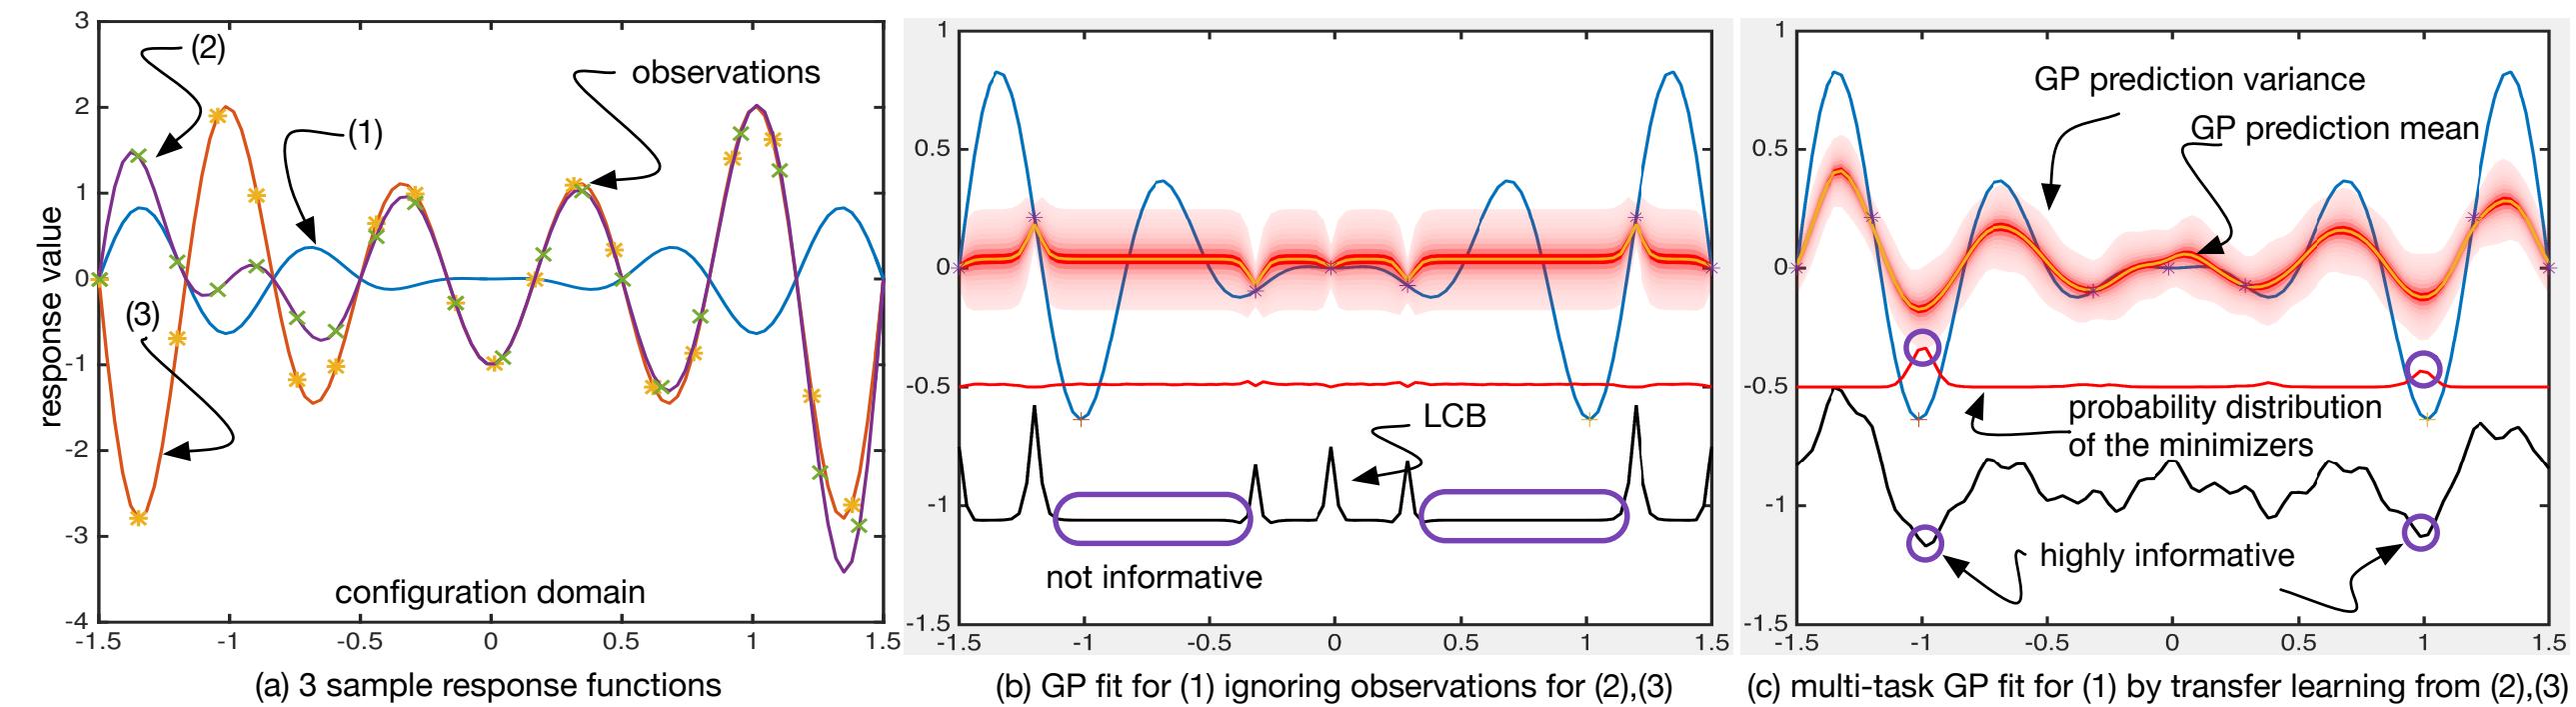
Algorithm 1 : TL4CO

Input: Configuration space \mathbb{X} , Number of historical configuration optimization datasets T , Maximum budget N_{max} , Response function f , Kernel function \mathbf{K} , Hyper-parameters $\boldsymbol{\theta}^j$, The current dataset $\mathbb{S}^{j=1}$, Datasets belonging to other versions $\mathbb{S}^{j=2:T}$, Diagonal noise matrix Σ , Design sample size n , Learning cycle N_l

Output: Optimal configurations \mathbf{x}^* and learned model \mathcal{M}

- 1: $\mathcal{D} = \{\mathbf{x}_1, \dots, \mathbf{x}_n\}$ create an initial sparse design (*lhd*)
 - 2: Obtain the *performance* $\forall x_i \in \mathcal{D}, y_i \leftarrow f(\mathbf{x}_i) + \epsilon_i$
 - 3: $\mathbb{S}^{j=2:T} \leftarrow$ select m points from other versions ($j = 2 : T$) that reduce entropy ▷ Eq.(8)
 - 4: $\mathbb{S}_{1:n}^1 \leftarrow \mathbb{S}^{j=2:T} \cup \{(\mathbf{x}_i, y_i)\}_{i=1}^n; t \leftarrow n + 1$
 - 5: $\mathcal{M}(\mathbf{x}|\mathbb{S}_{1:n}^1, \boldsymbol{\theta}^j) \leftarrow$ fit a multi-task \mathcal{GP} to \mathcal{D} ▷ Eq. (6)
 - 6: **while** $t \leq N_{max}$ **do**
 - 7: If $(t \bmod N_l = 0)$ then $[\boldsymbol{\theta} \leftarrow$ learn the kernel hyper-parameters by maximizing the likelihood]
 - 8: Determine the *next configuration*, \mathbf{x}_t , by optimizing LCB over \mathcal{M} , $\mathbf{x}_t \leftarrow \arg \max_{\mathbf{x}} u(\mathbf{x}|\mathcal{M}, \mathbb{S}_{1:t-1}^1)$ ▷ Eq.(11)
 - 9: Obtain the *performance* of \mathbf{x}_t , $y_t \leftarrow f(\mathbf{x}_t) + \epsilon_t$
 - 10: Augment the configuration $\mathbb{S}_{1:t}^1 = \mathbb{S}_{1:t-1}^1 \cup \{(\mathbf{x}_t, y_t)\}$
 - 11: $\mathcal{M}(\mathbf{x}|\mathbb{S}_{1:t}^1, \boldsymbol{\theta}) \leftarrow$ re-fit a new GP model ▷ Eq.(6)
 - 12: $t \leftarrow t + 1$
 - 13: **end while**
 - 14: $(\mathbf{x}^*, y^*) = \min[\mathbb{S}_{1:N_{max}}^1 - \mathbb{S}^{j=2:T}]$ locating the optimum configuration by discarding data for other system versions
 - 15: $\mathcal{M}(\mathbf{x})$ storing the learned model
-

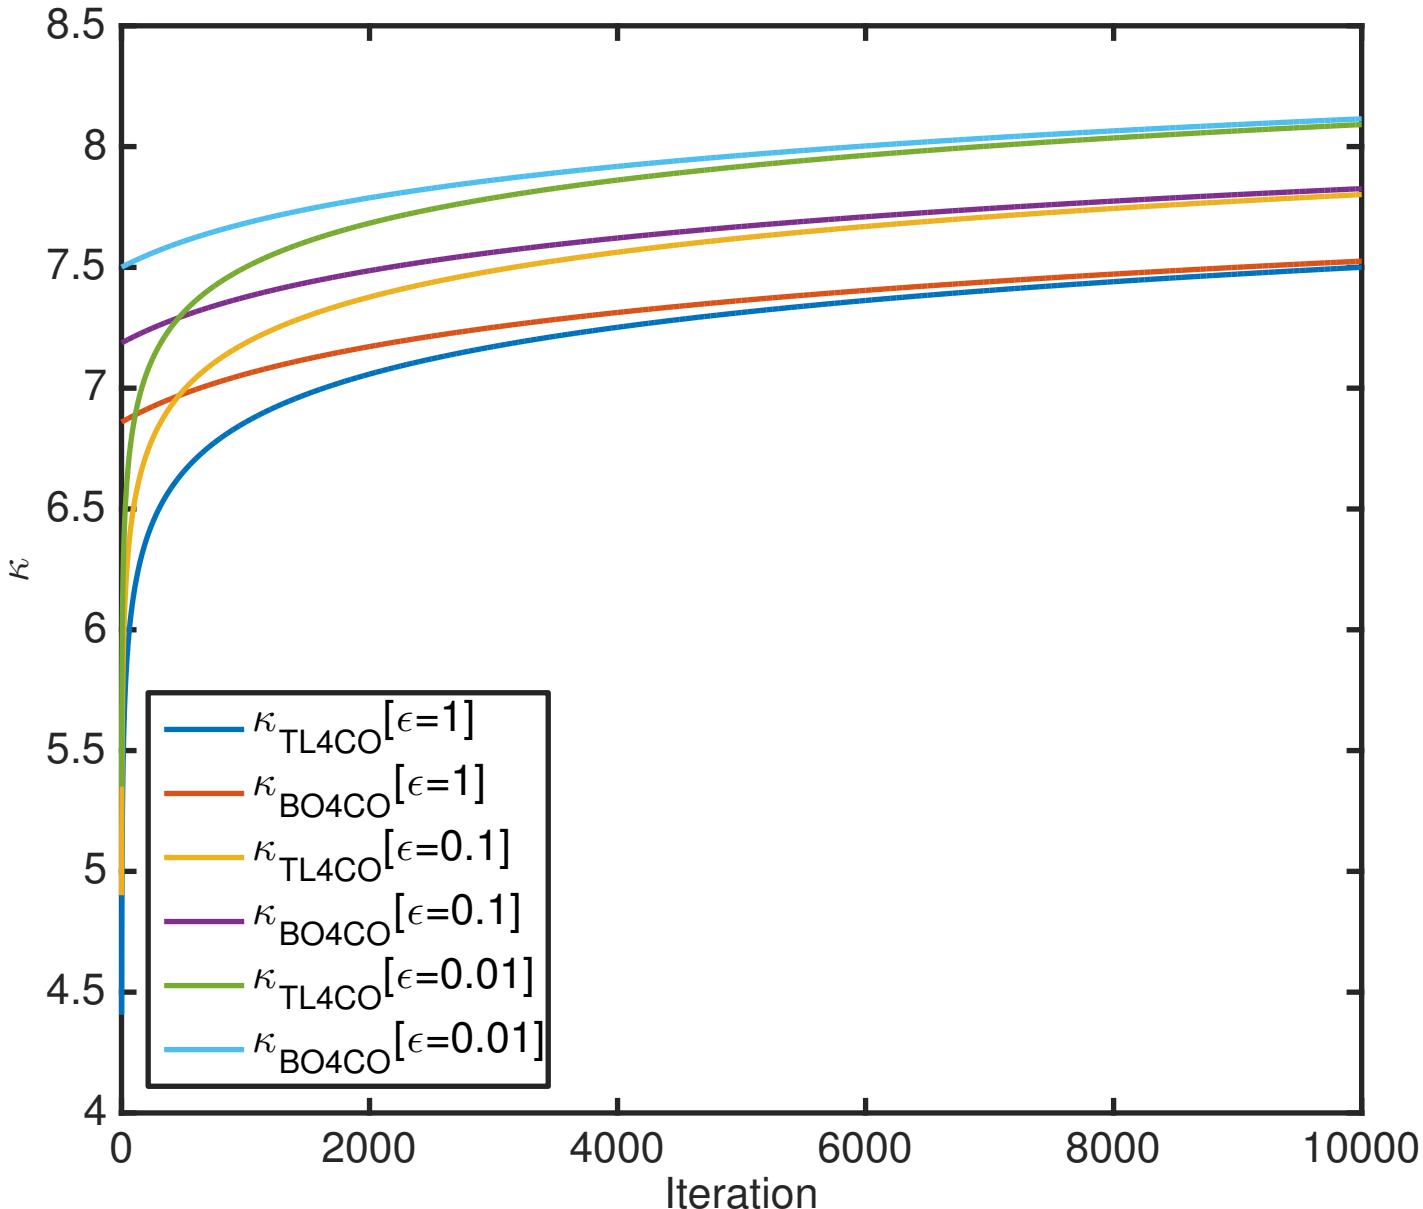
Multi-task GP vs single-task GP



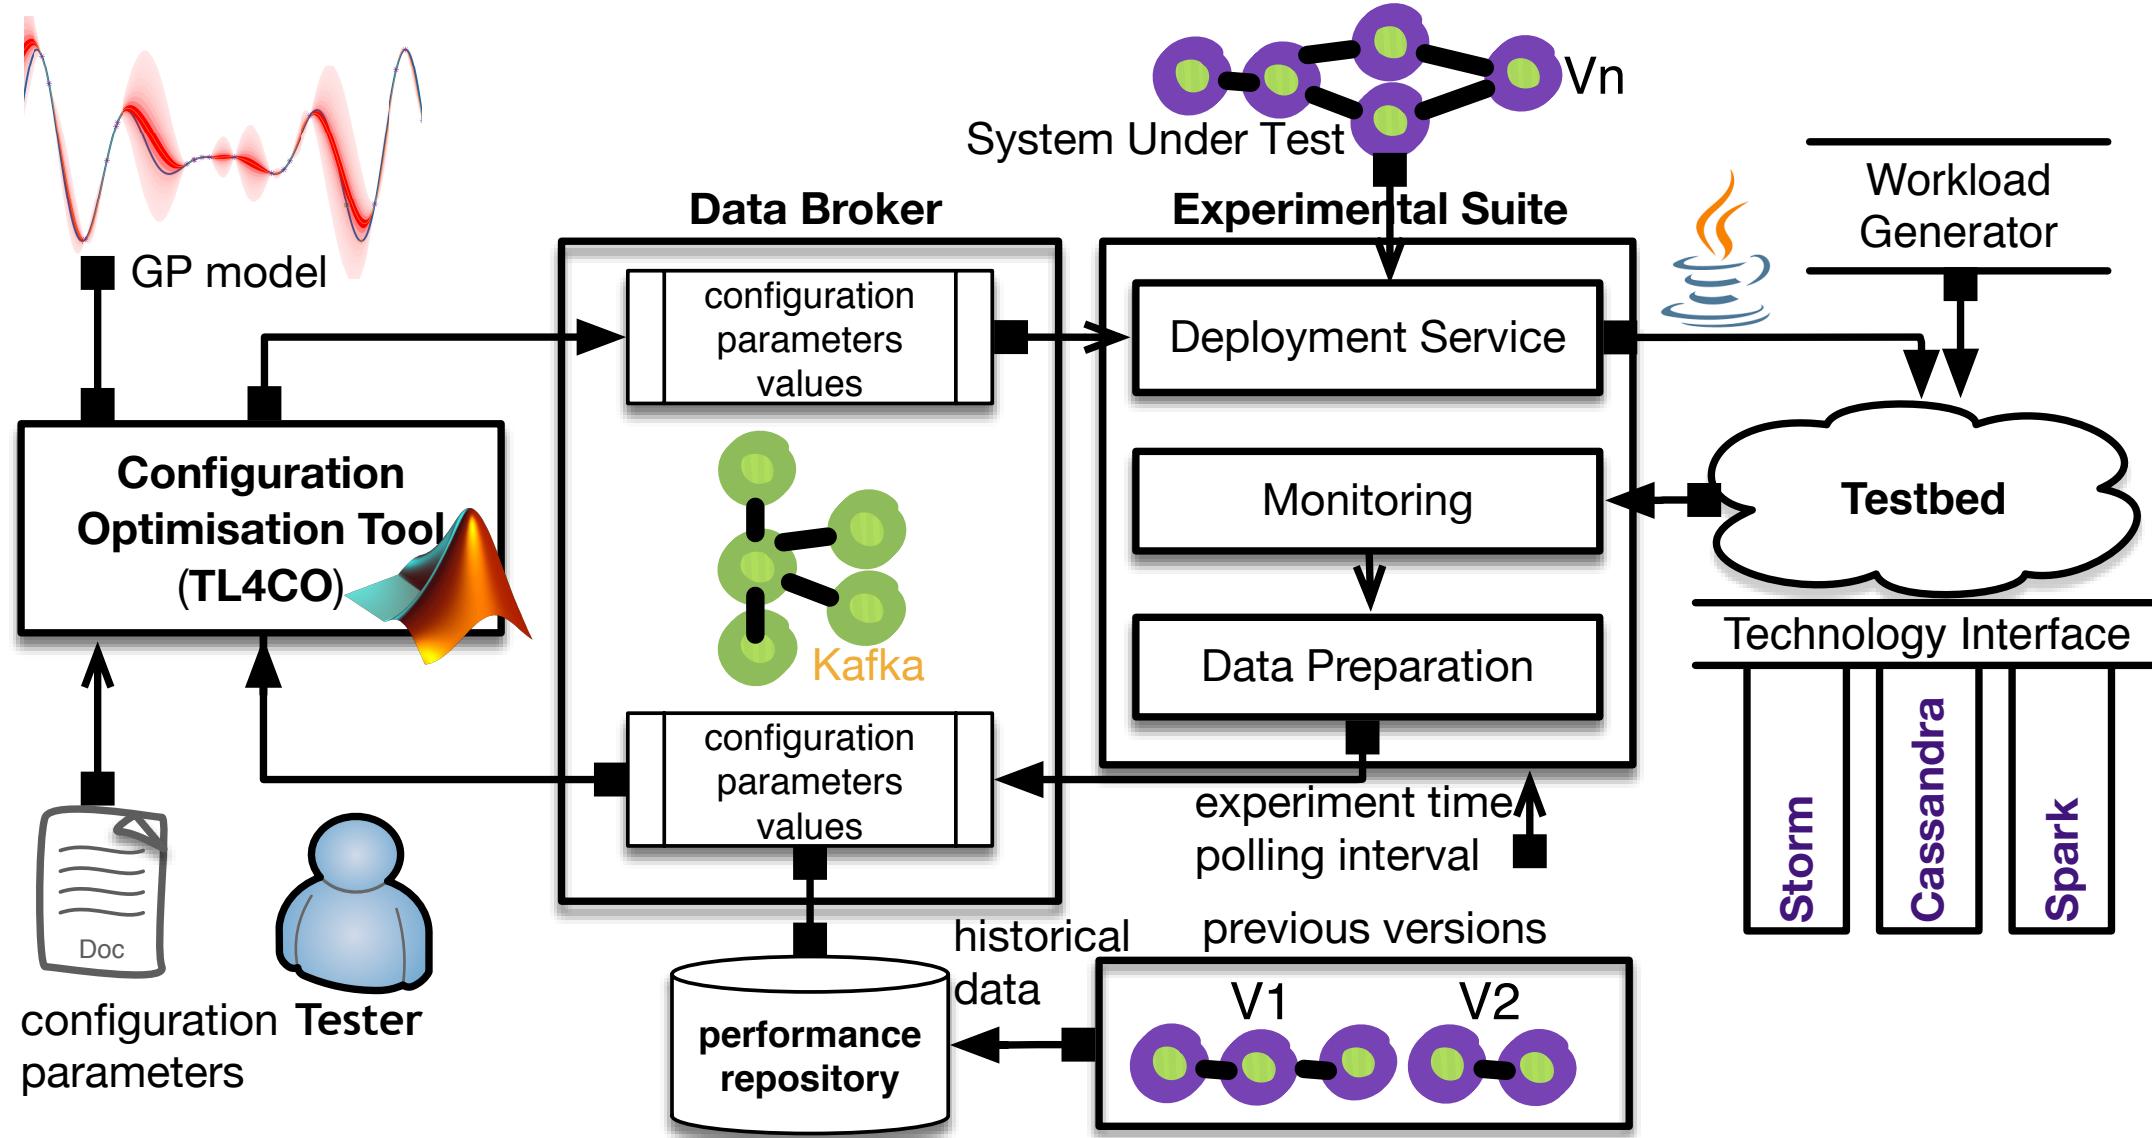
Exploitation vs exploration

$$\boldsymbol{x}_{t+1} = \operatorname{argmax}_{\boldsymbol{x} \in \mathbb{X}} u(\boldsymbol{x} | \mathcal{M}, \mathbb{S}_{1:t}^1)$$

$$u_{LCB}(\boldsymbol{x} | \mathcal{M}, \mathbb{S}_{1:t}^1) = \operatorname{argmin}_{\boldsymbol{x} \in \mathbb{X}} \mu_t(\boldsymbol{x}) - \kappa \sigma_t(\boldsymbol{x}),$$

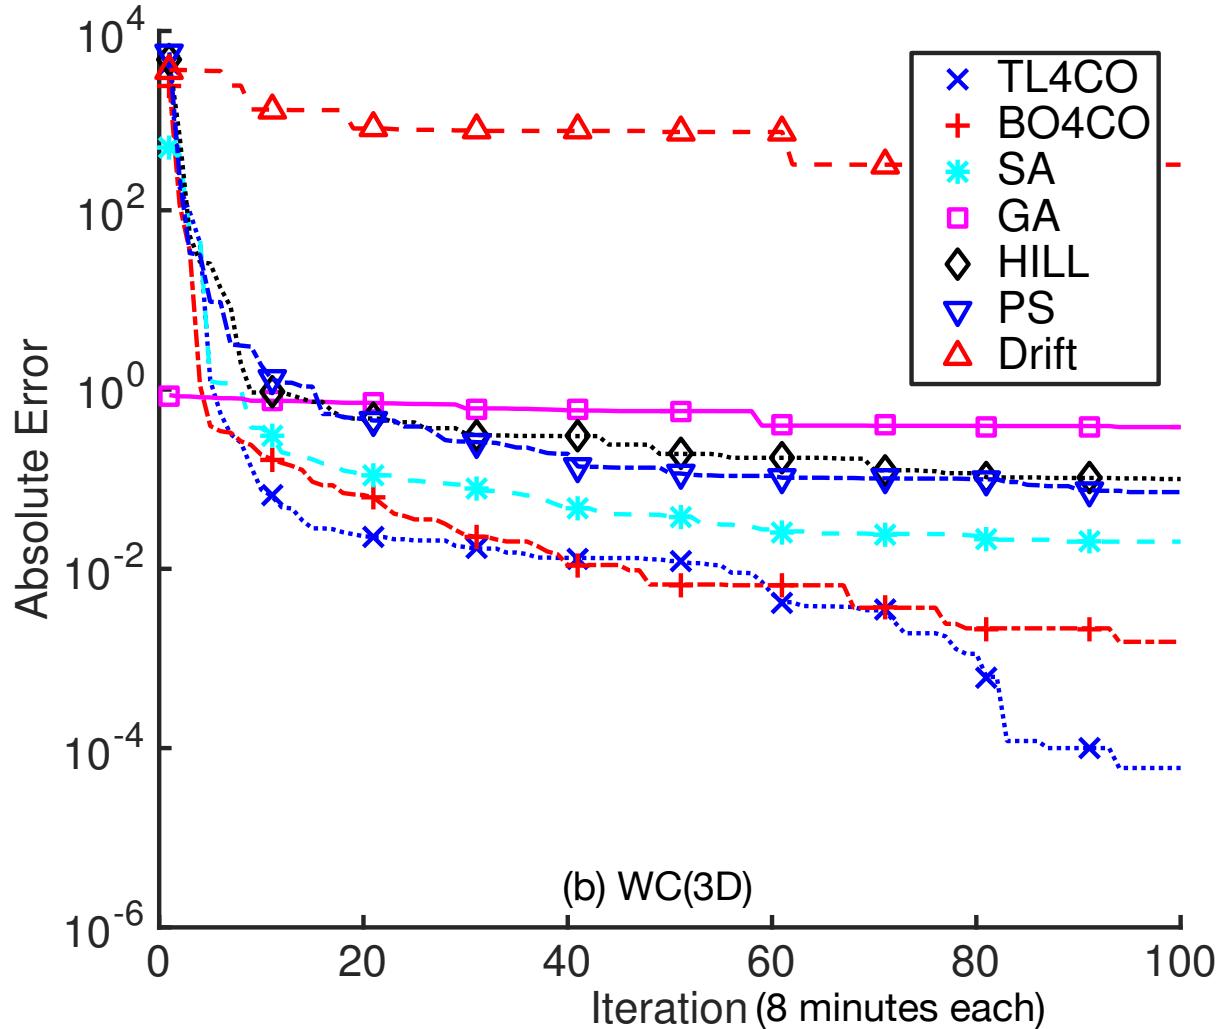


TL4CO architecture

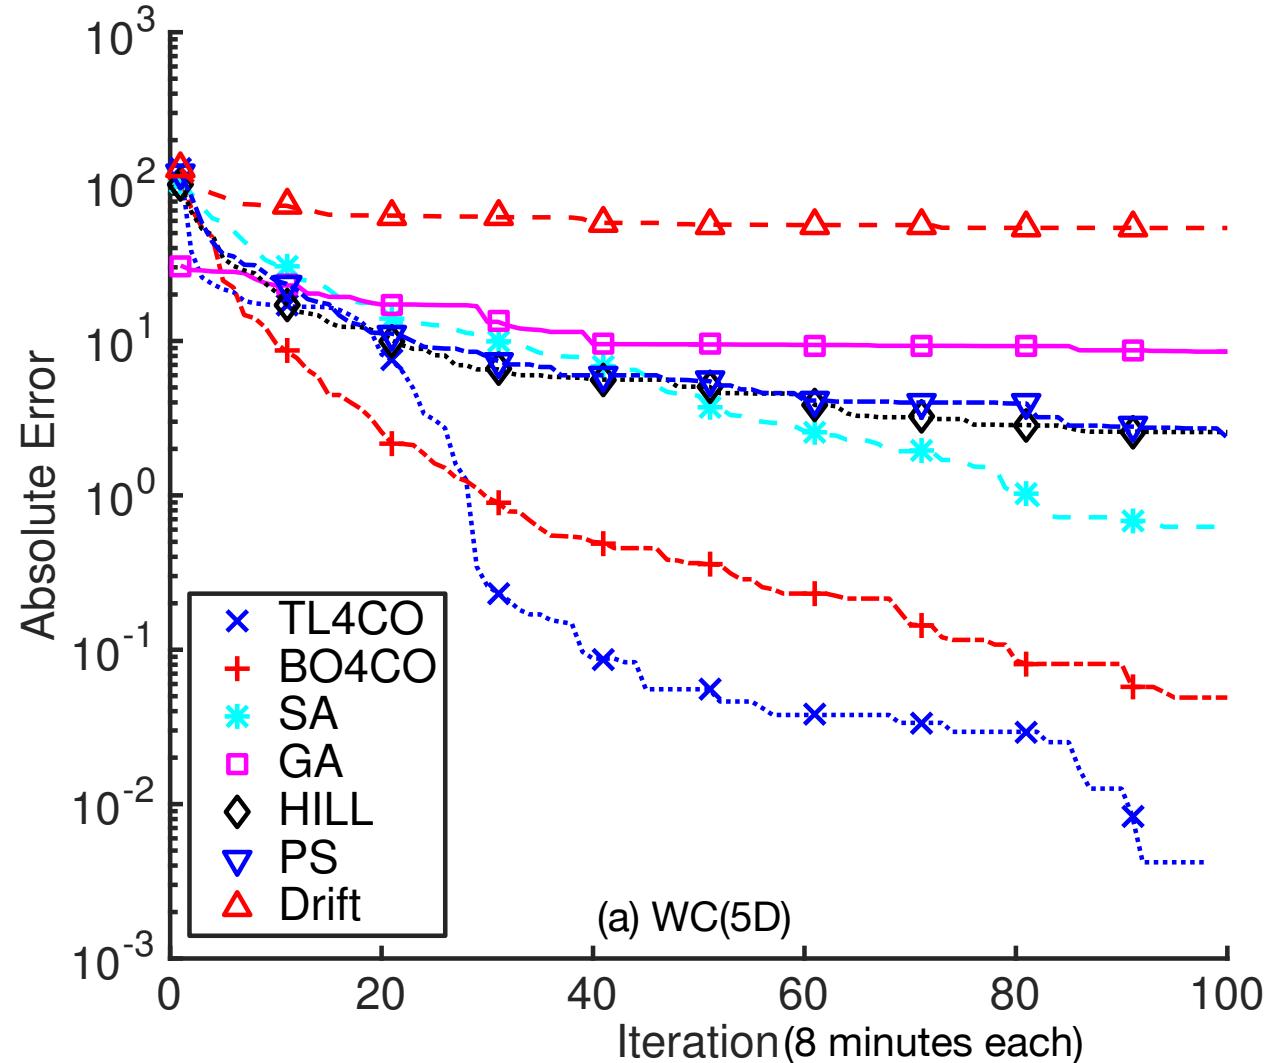


Experimental results

- 30 runs, report average performance
- Yes, we did full factorial measurements and we know where global min is... ☺

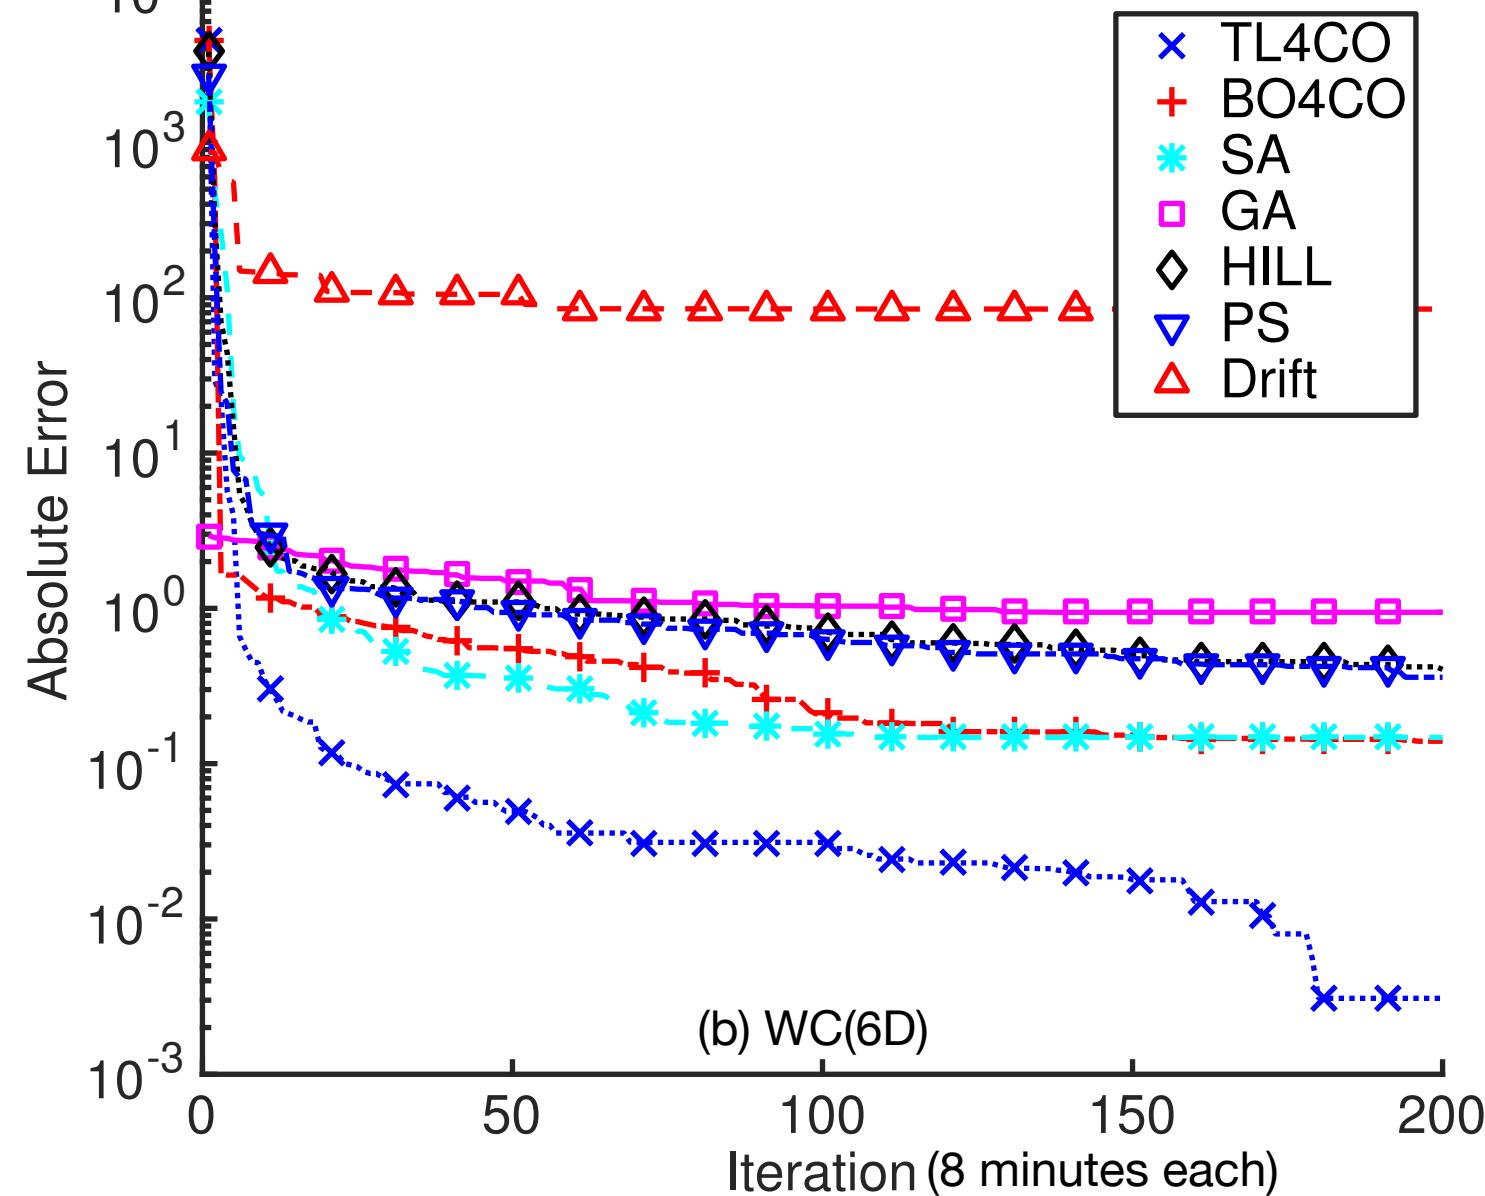


(b) WC(3D)

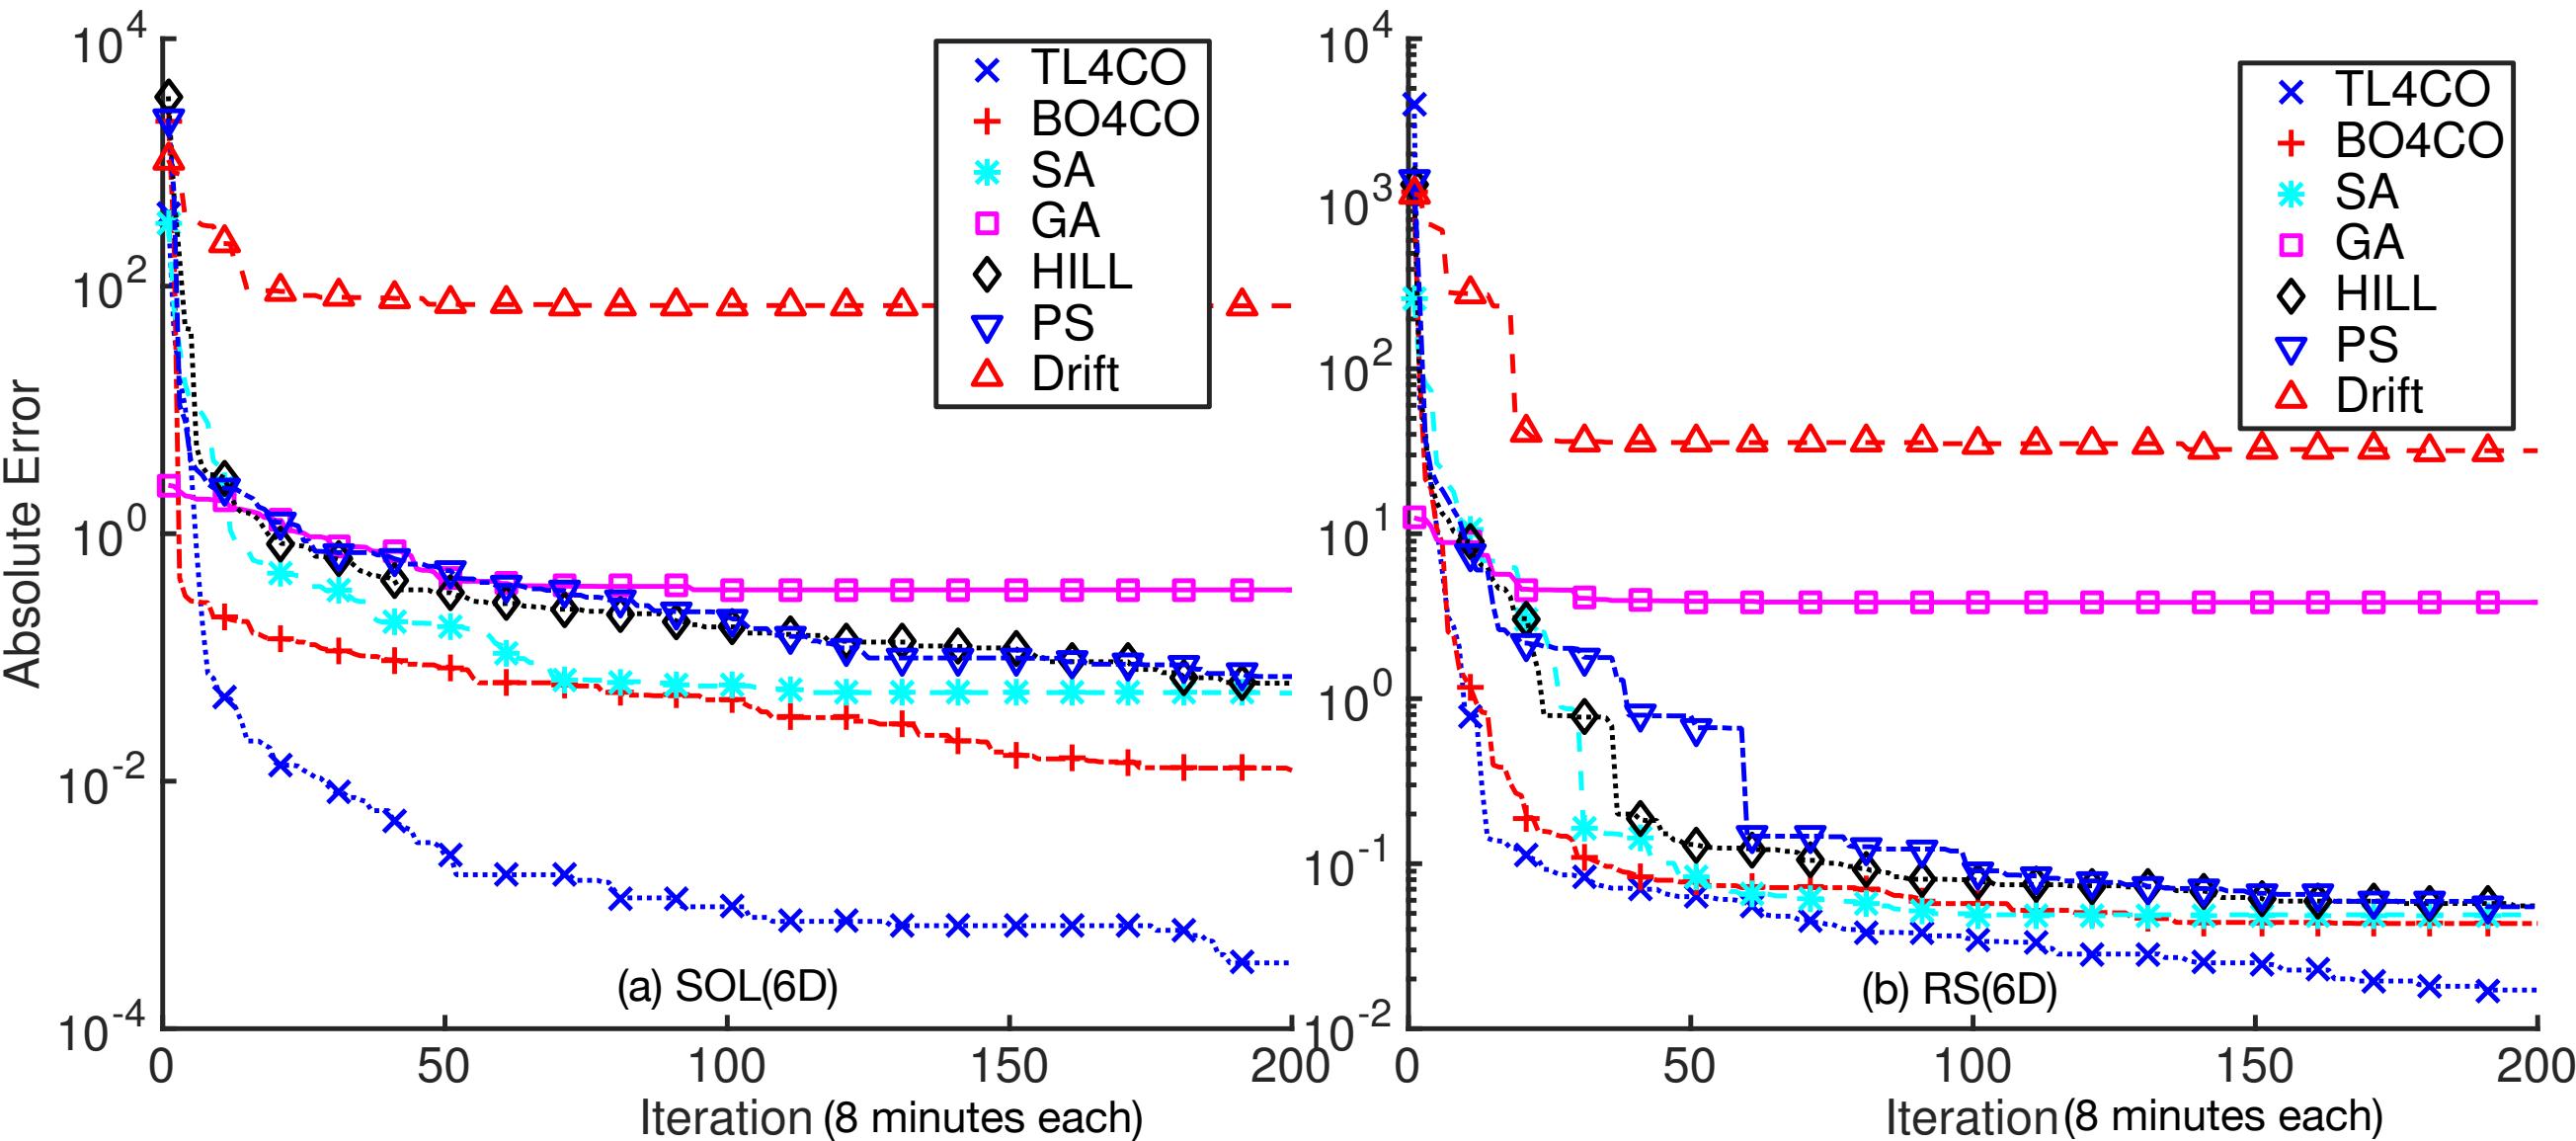


(a) WC(5D)

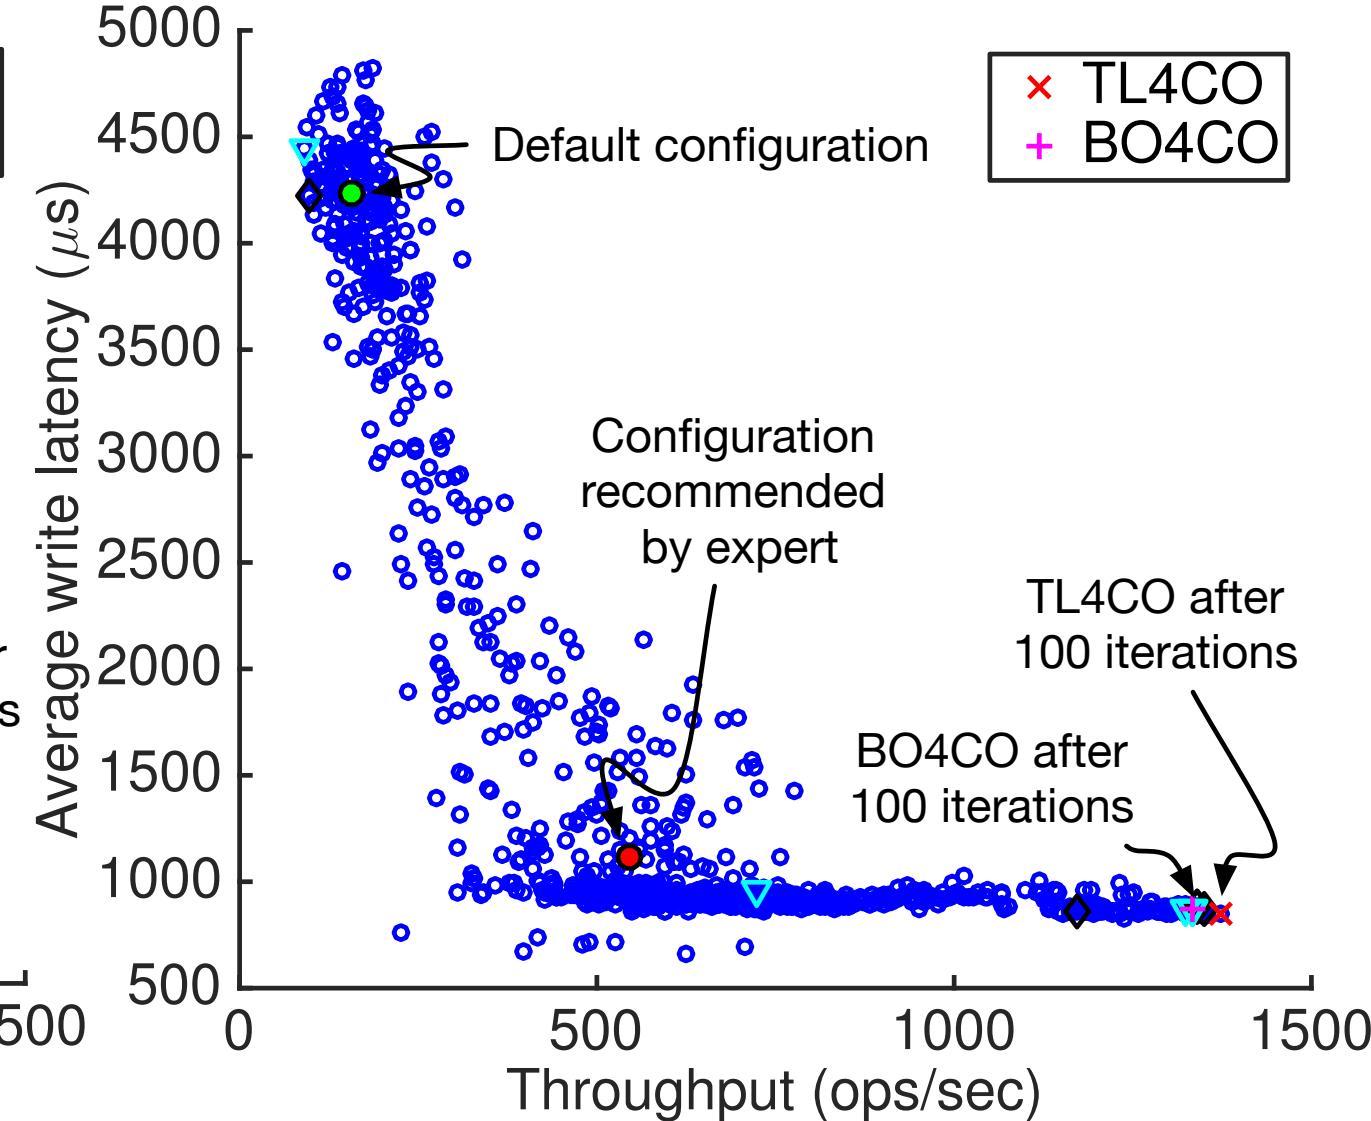
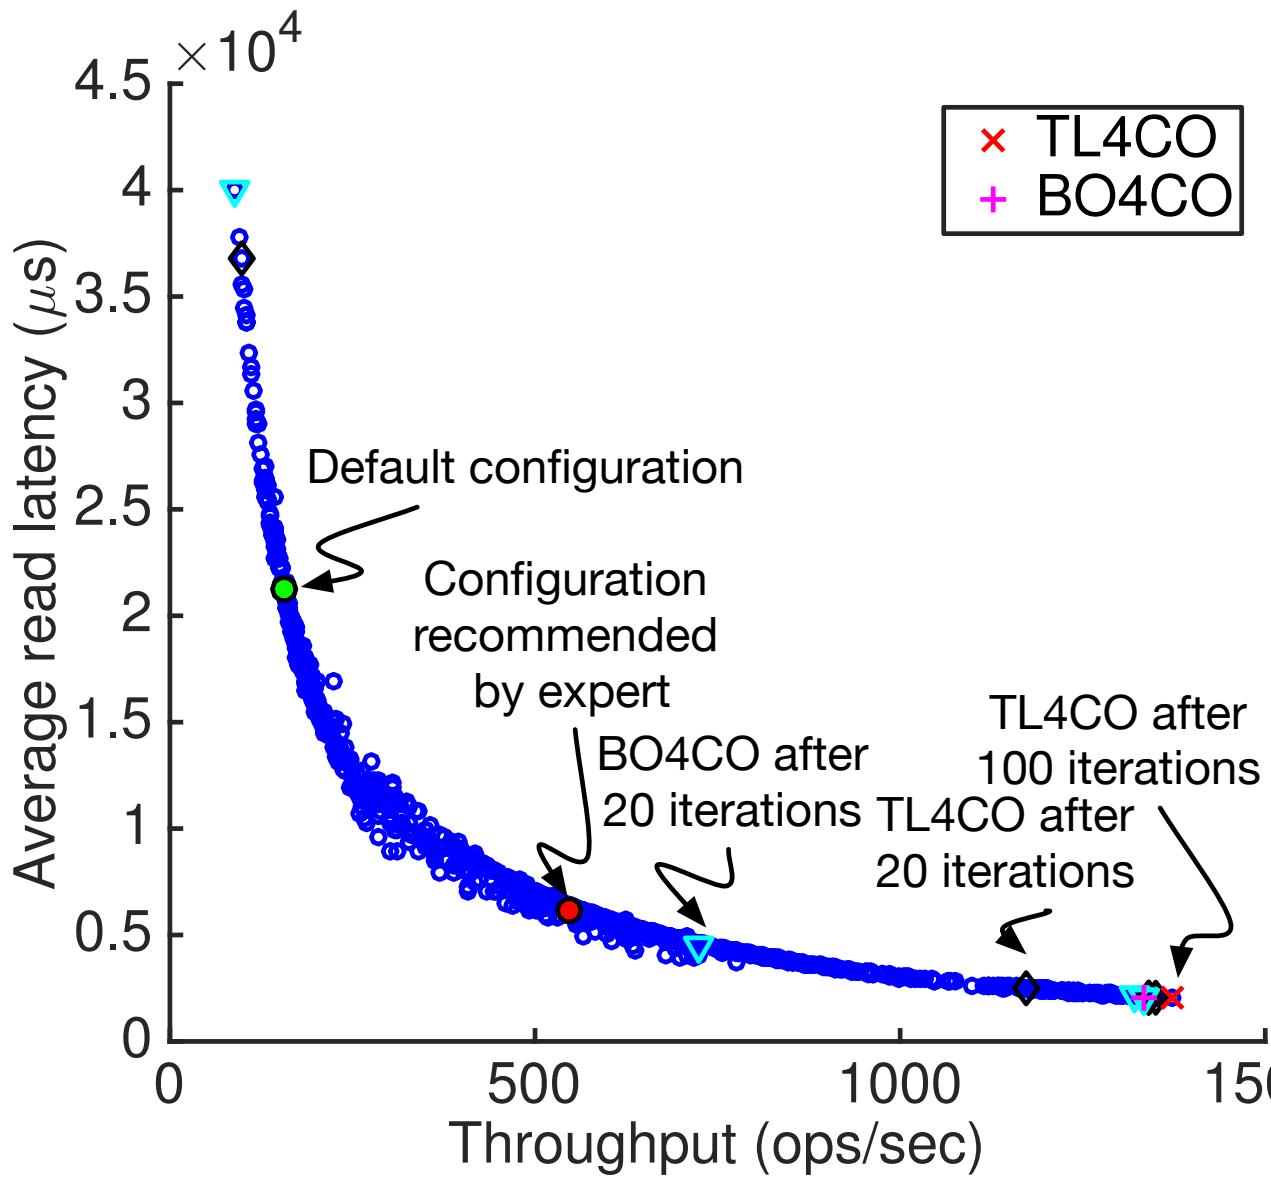
Experimental results (cont.)



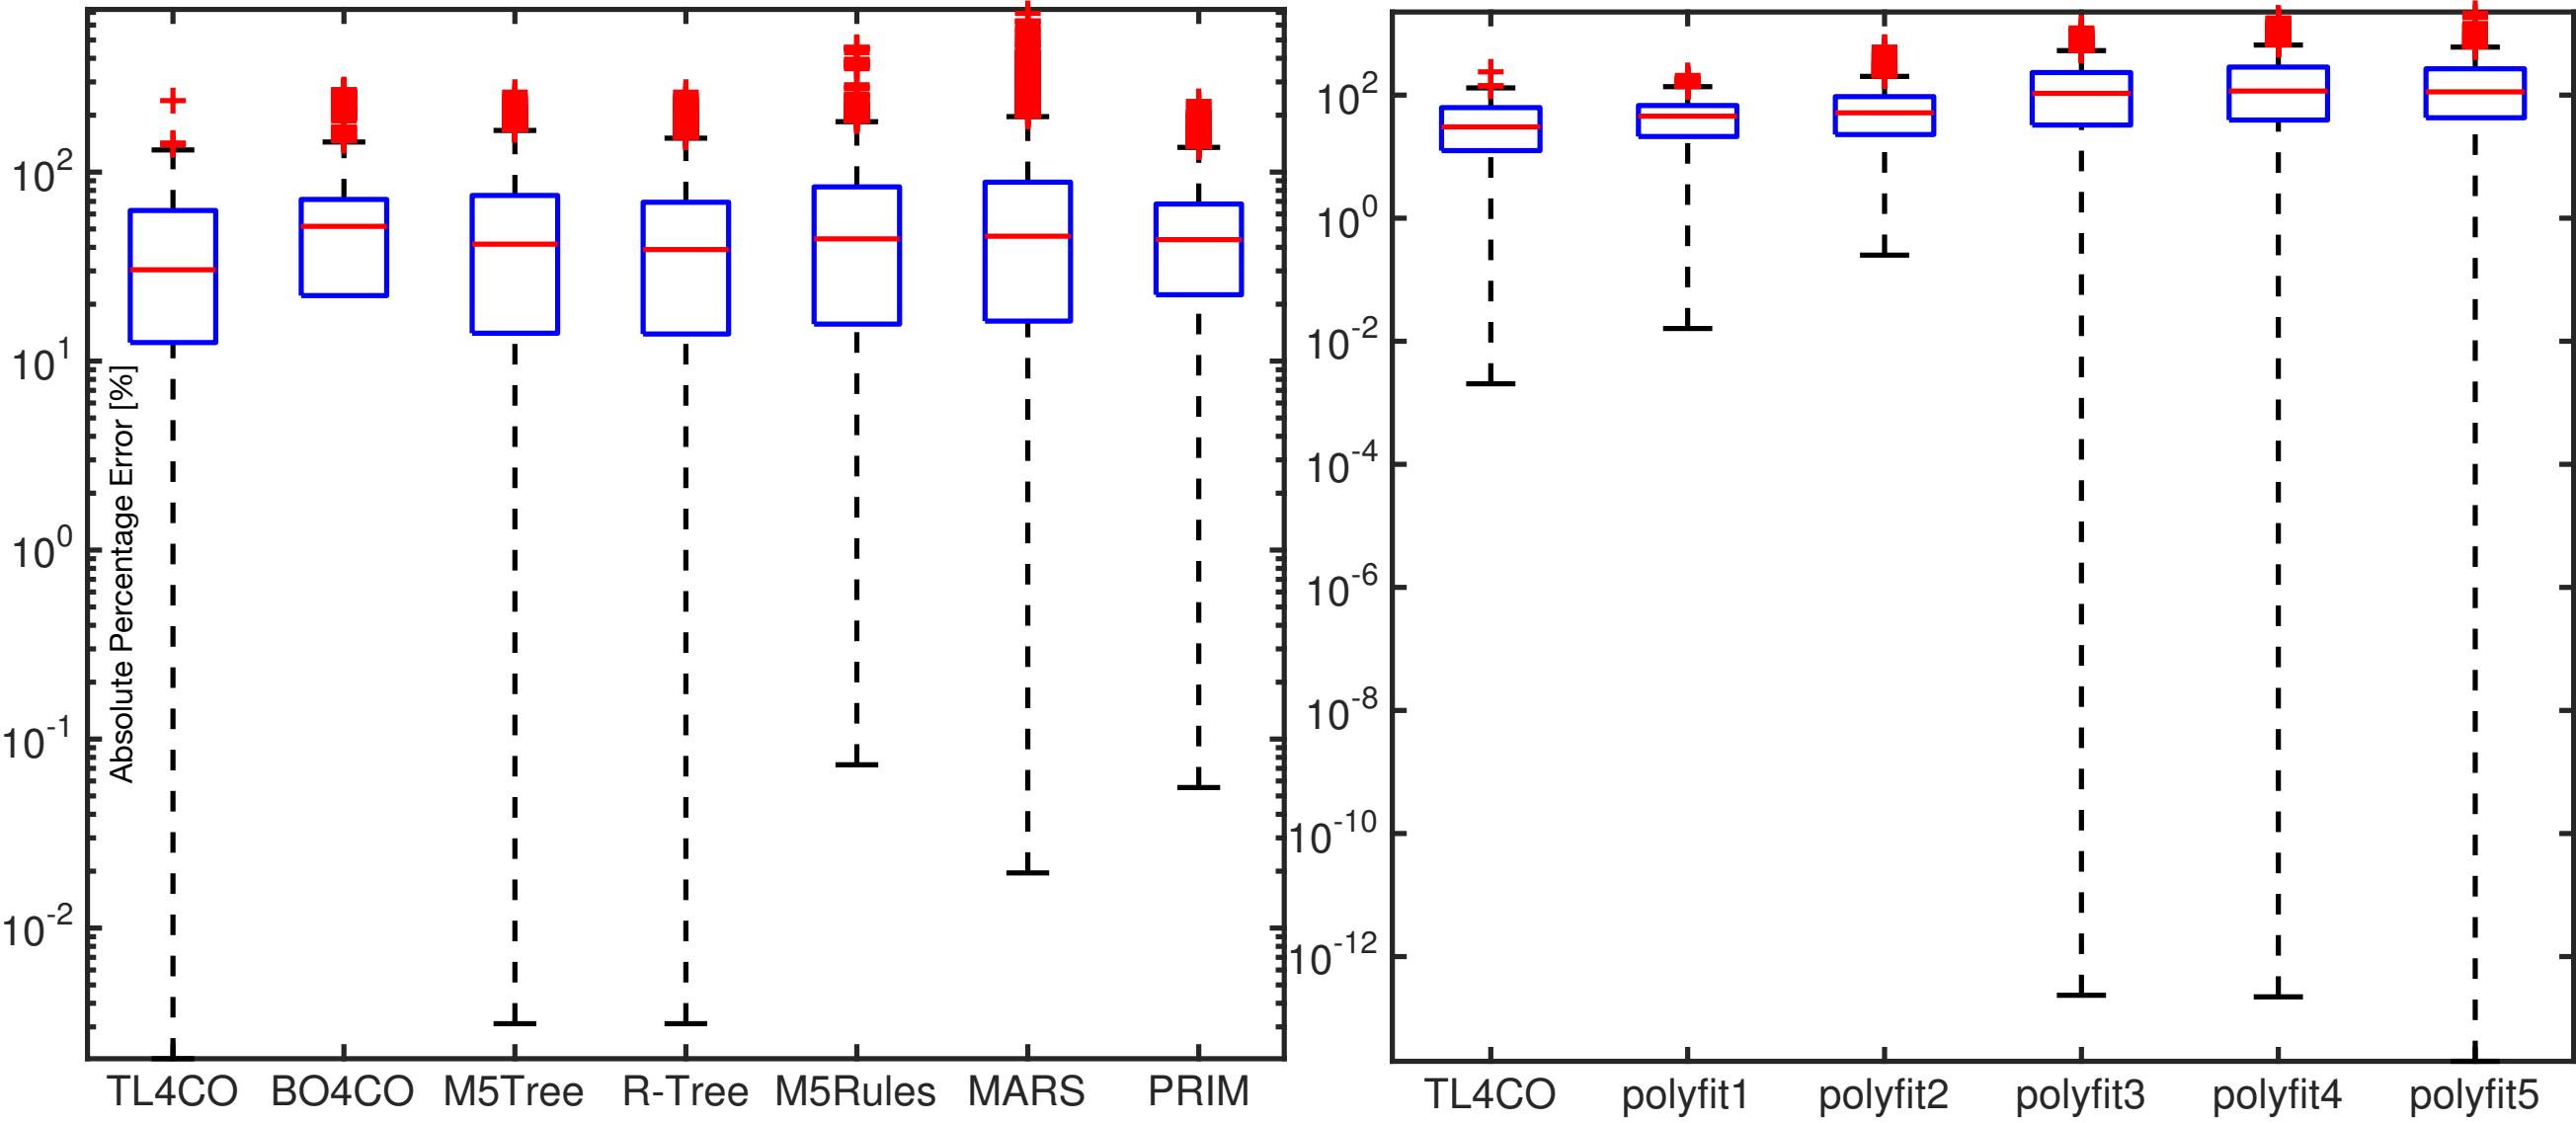
Experimental results (cont.)



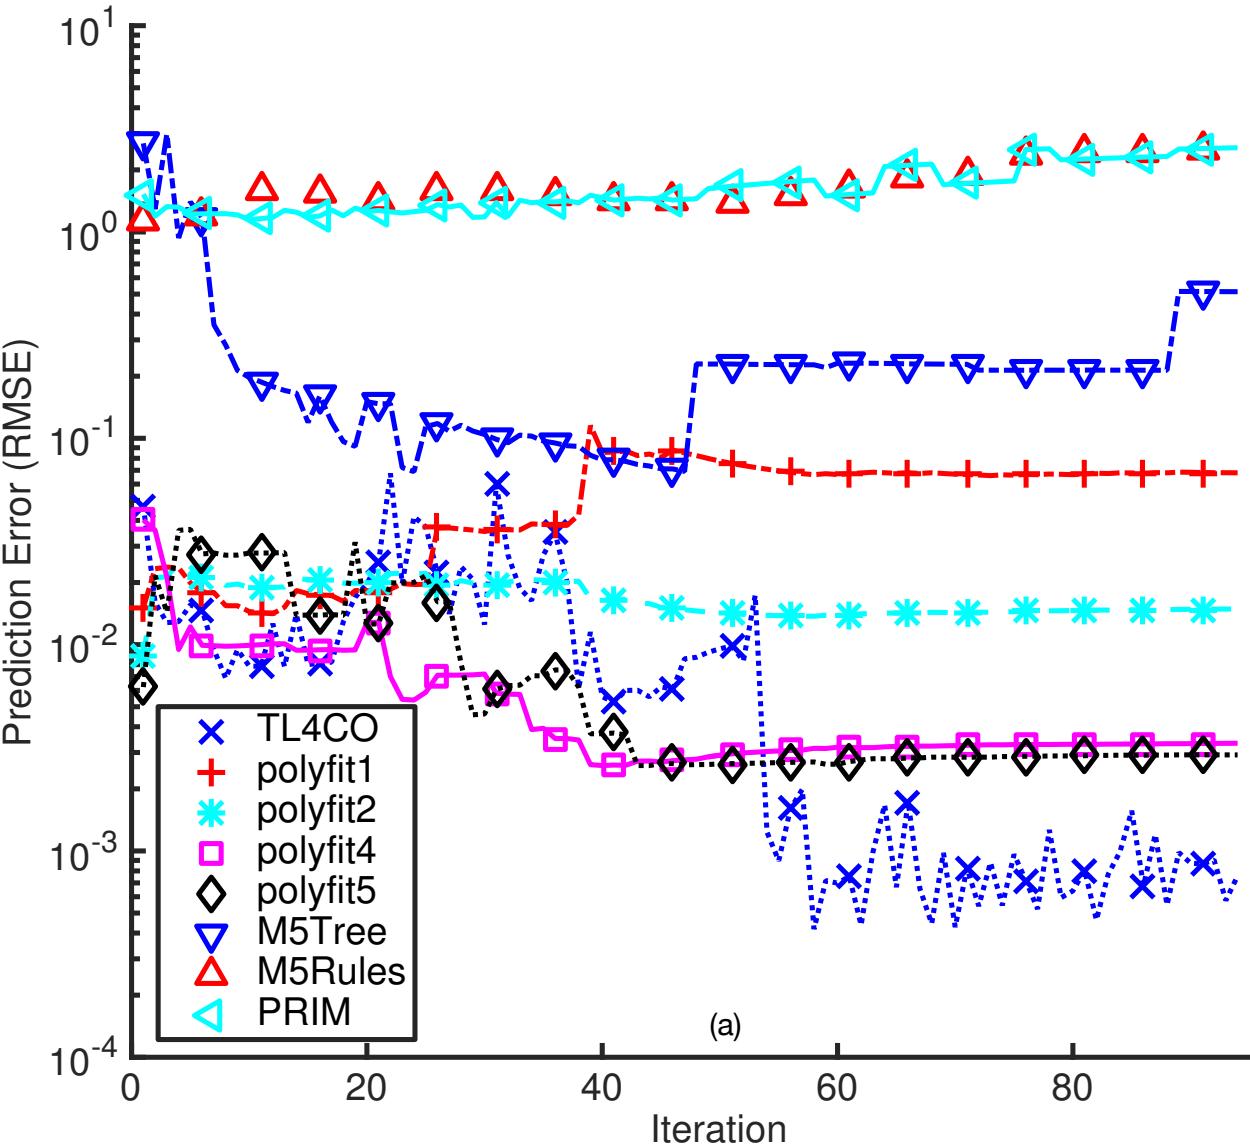
Comparison with default and expert prescription



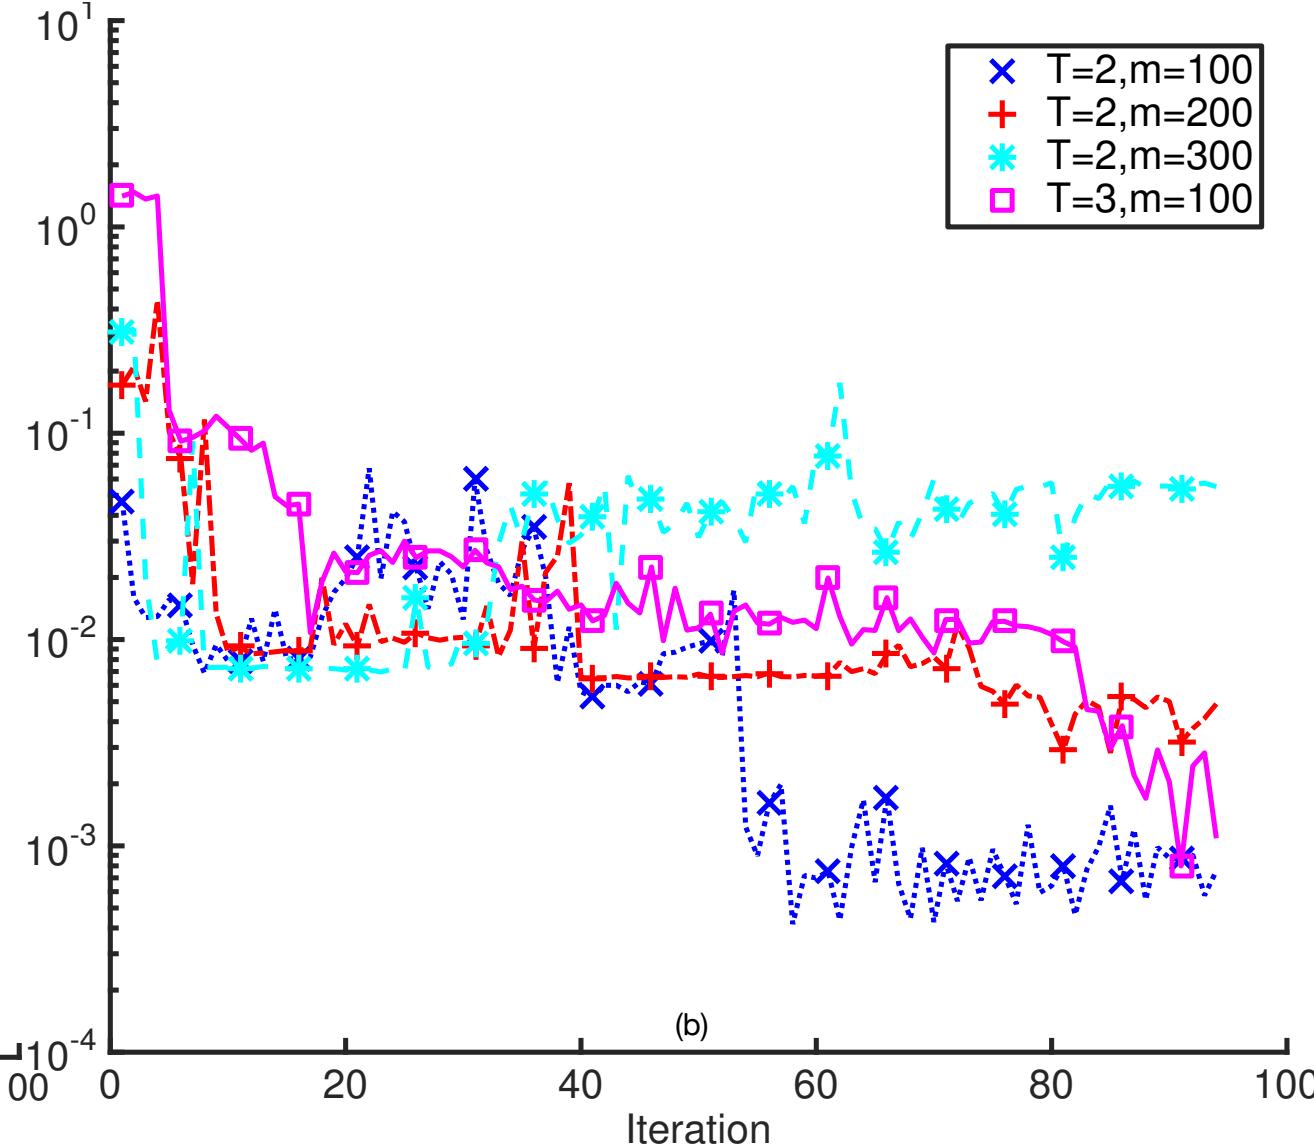
Model accuracy



Prediction accuracy over time



(a)

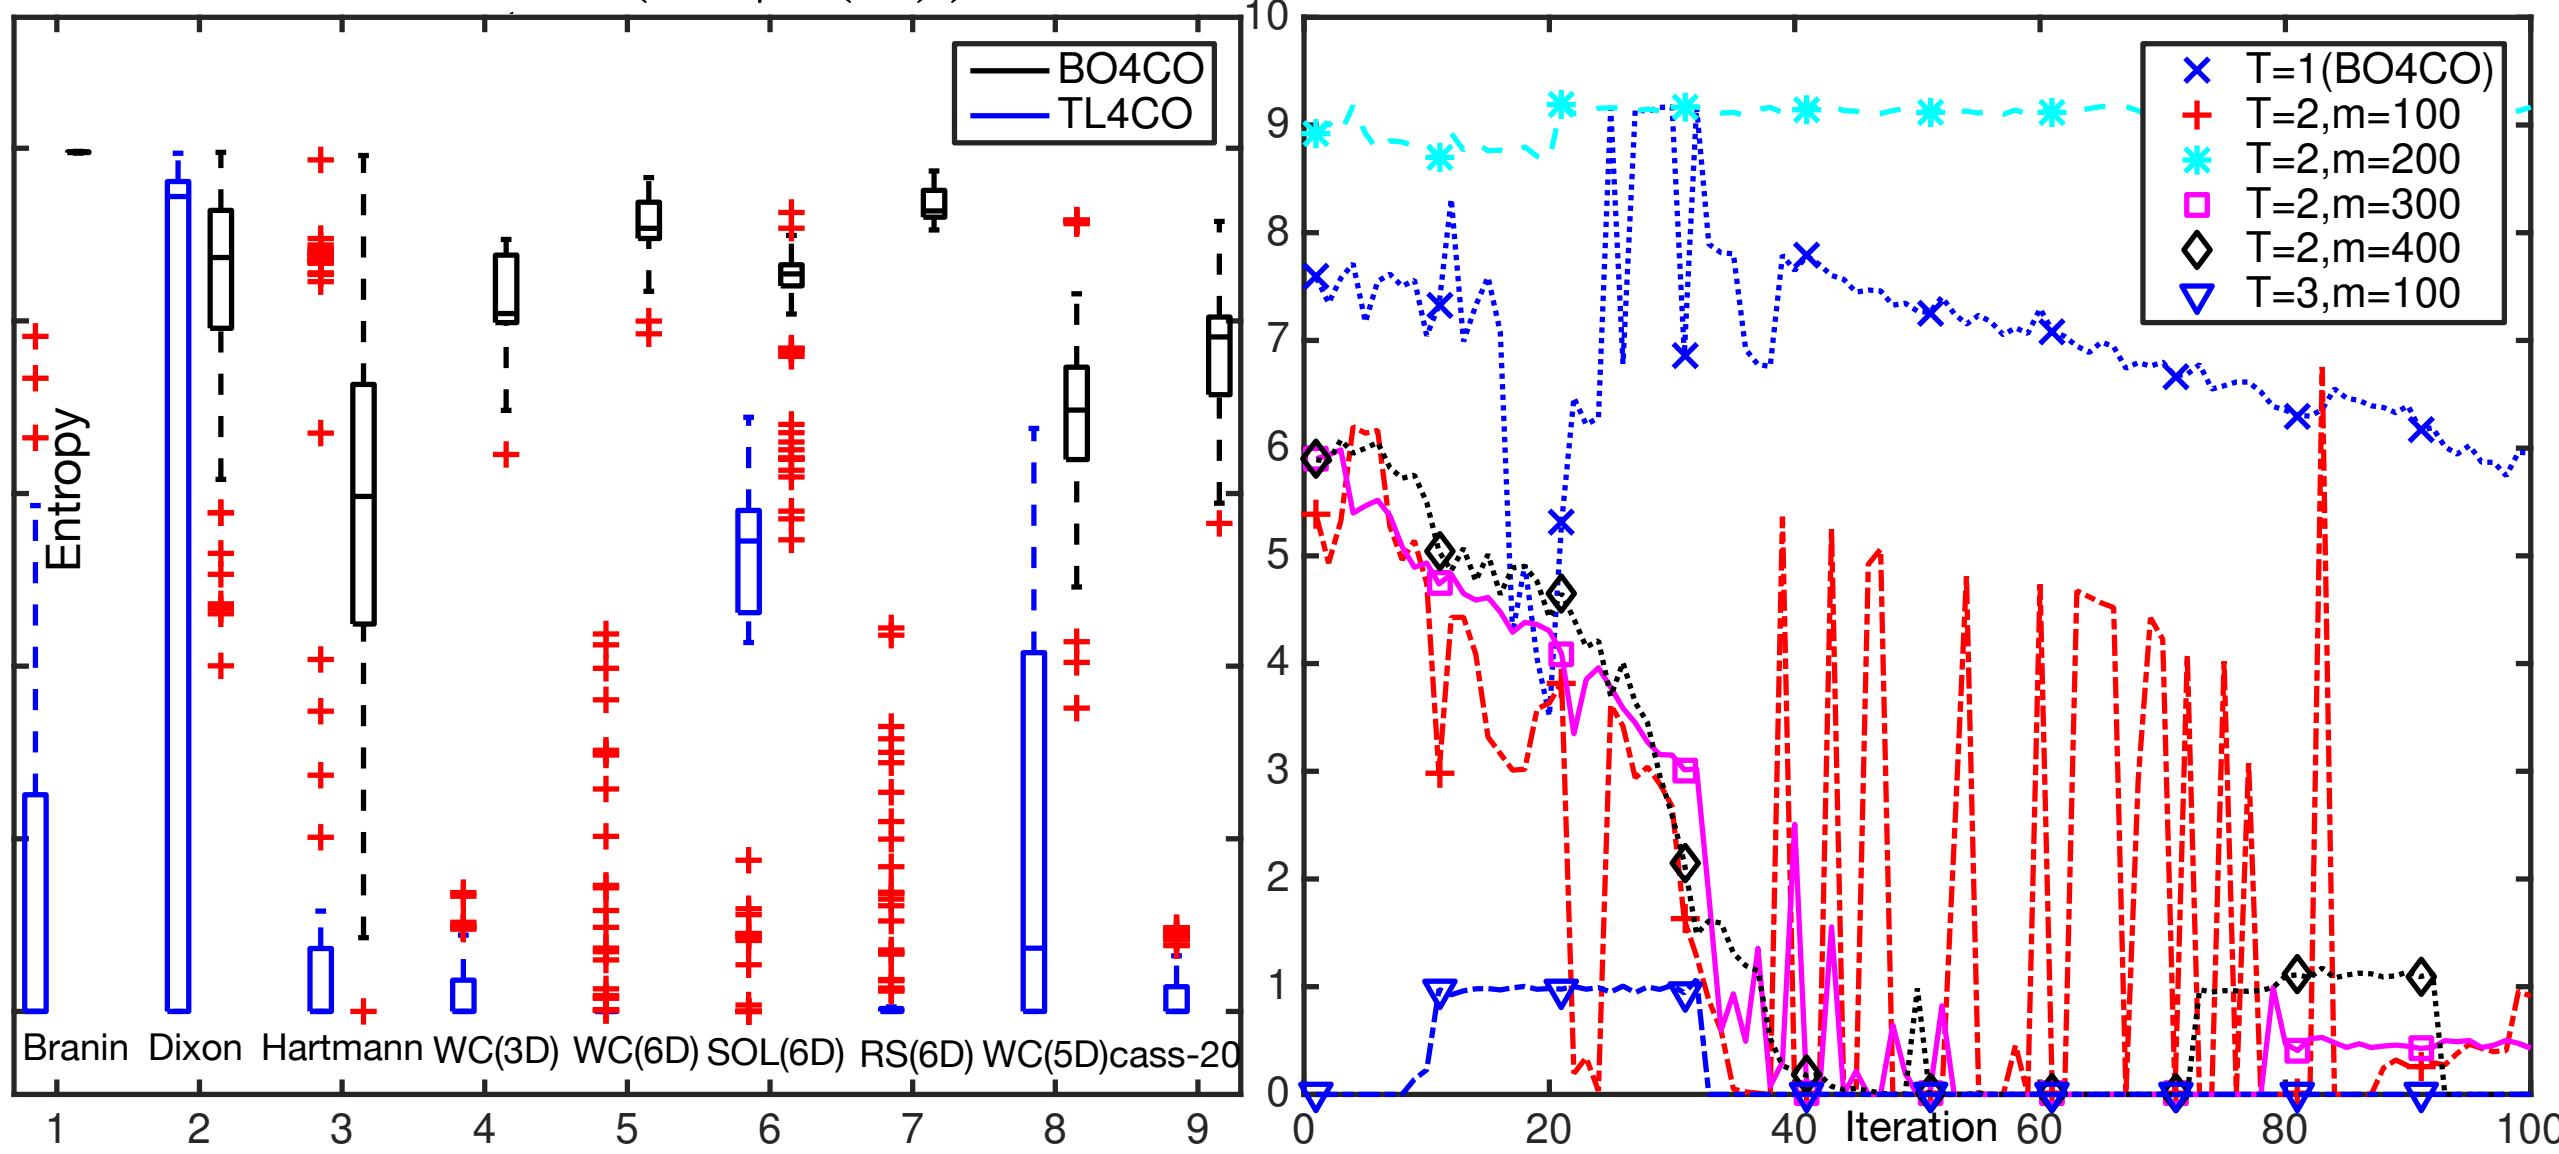


(b)

Entropy of the density function of the minimizers

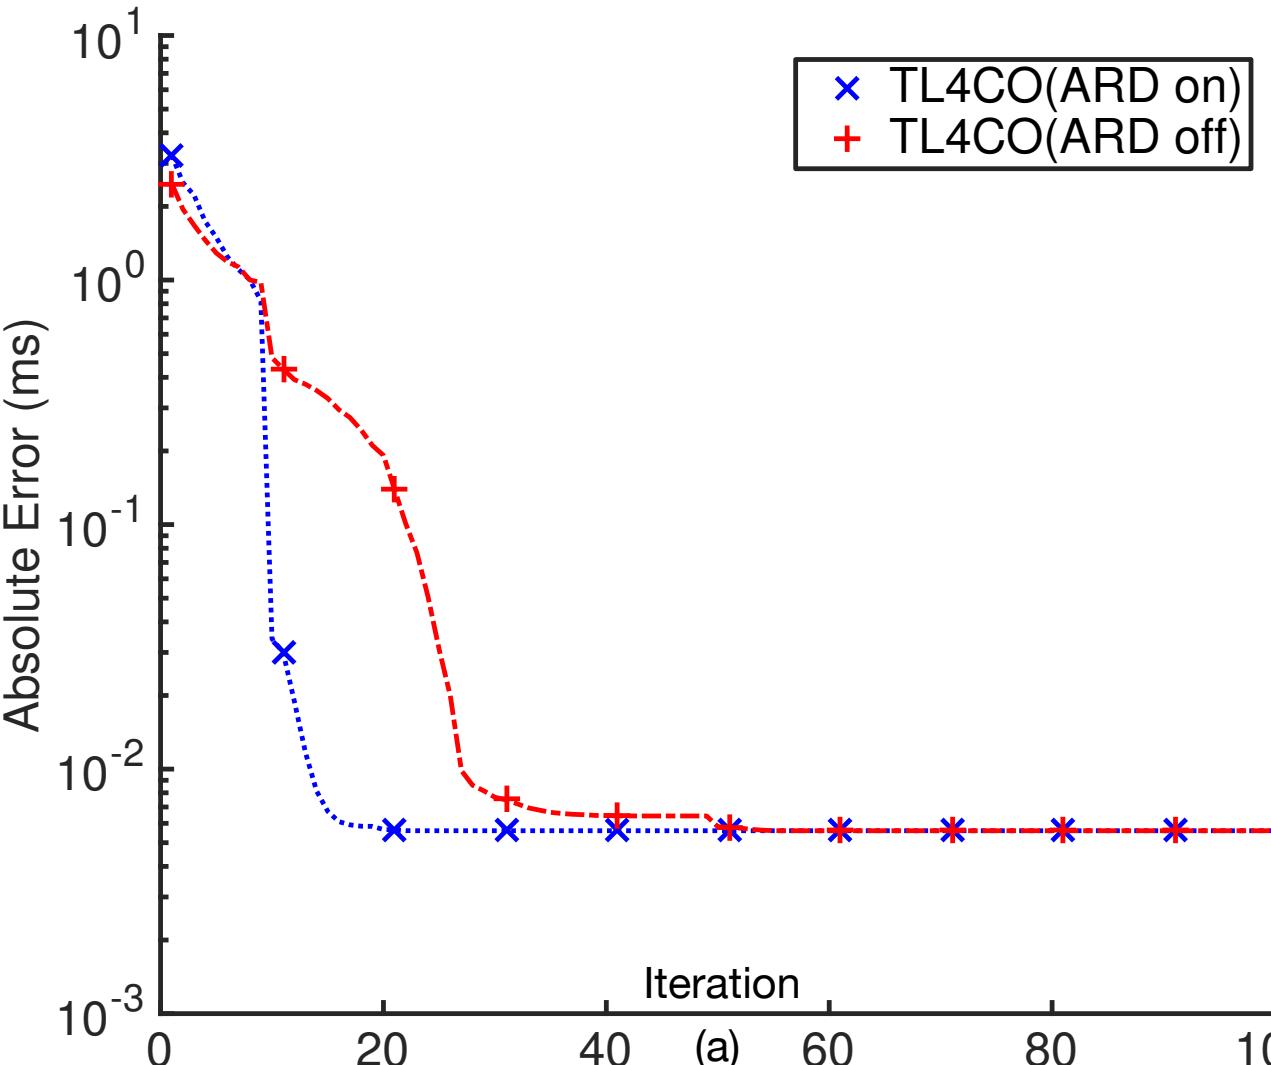
$$X^* = \Pr(\mathbf{x}^* | f(\mathbf{x}))$$

Knowledge about the location of the minimizer

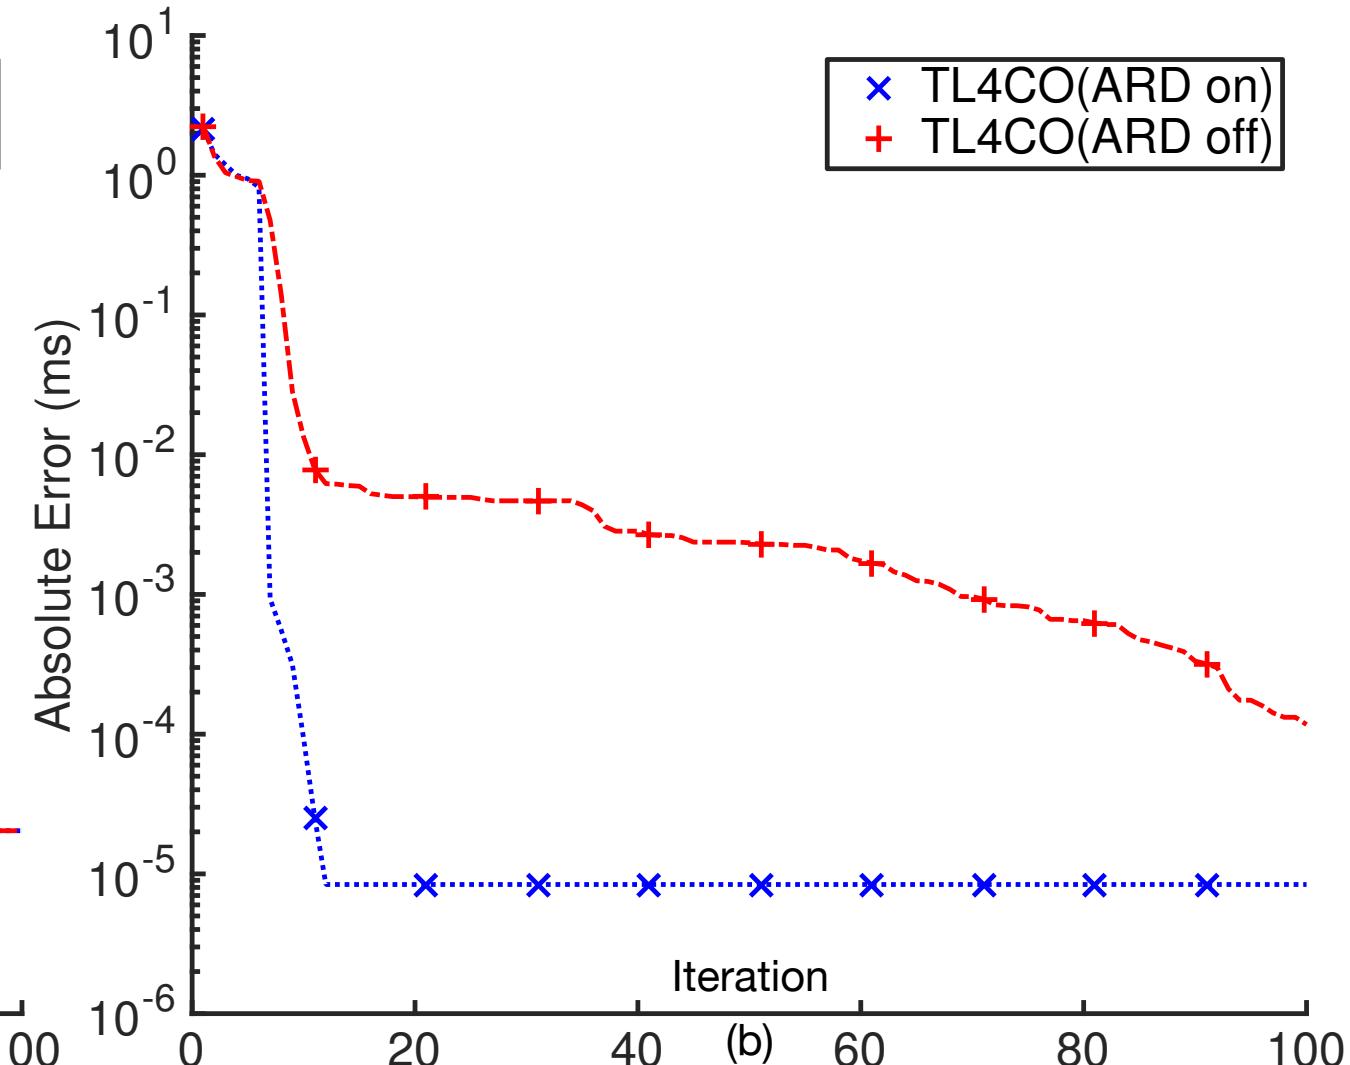


Effect of ARD

$$k_{xx}(\mathbf{x}_i, \mathbf{x}_j) = \exp\left(\sum_{\ell=1}^d (-\theta_\ell \delta(\mathbf{x}_i \neq \mathbf{x}_j))\right),$$

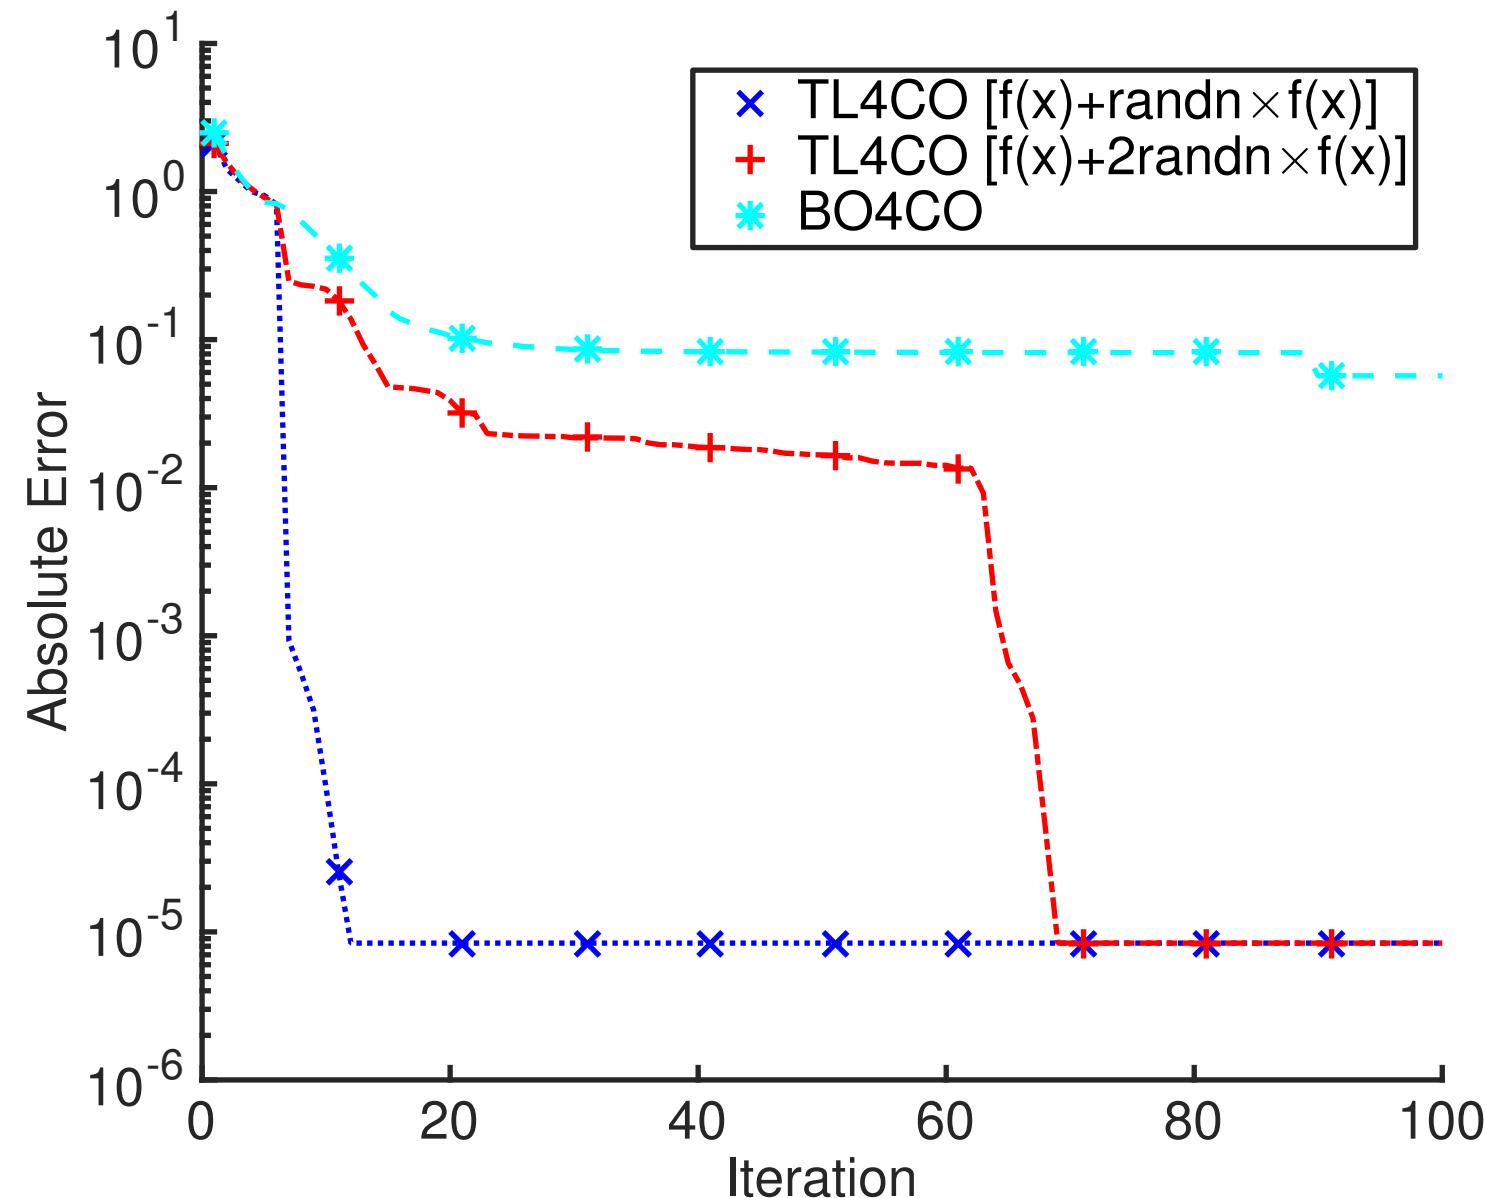


(a)

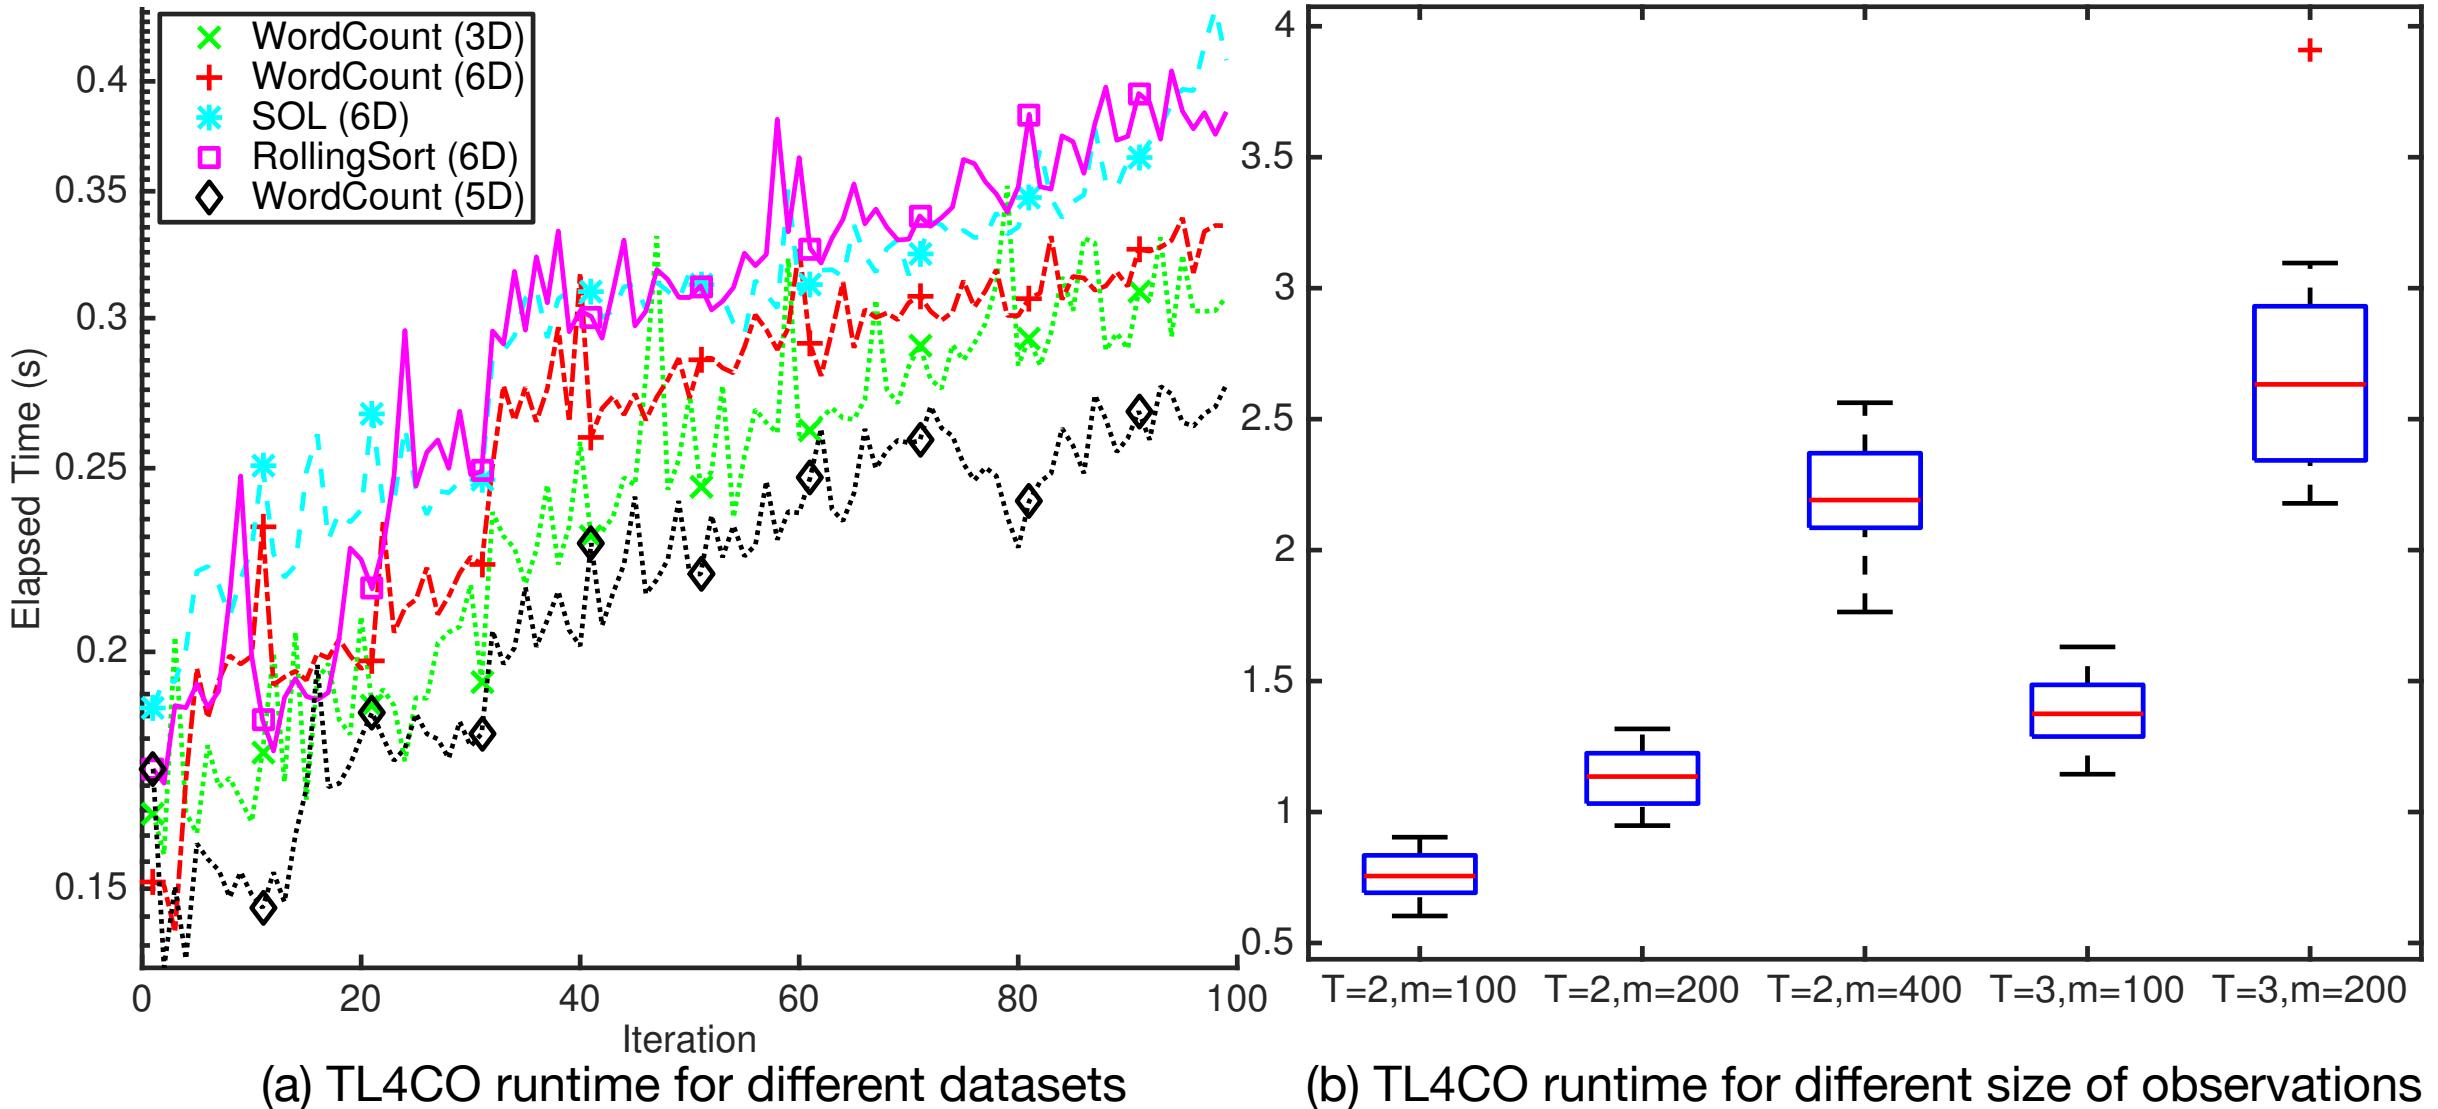


(b)

Effect of correlation between tasks



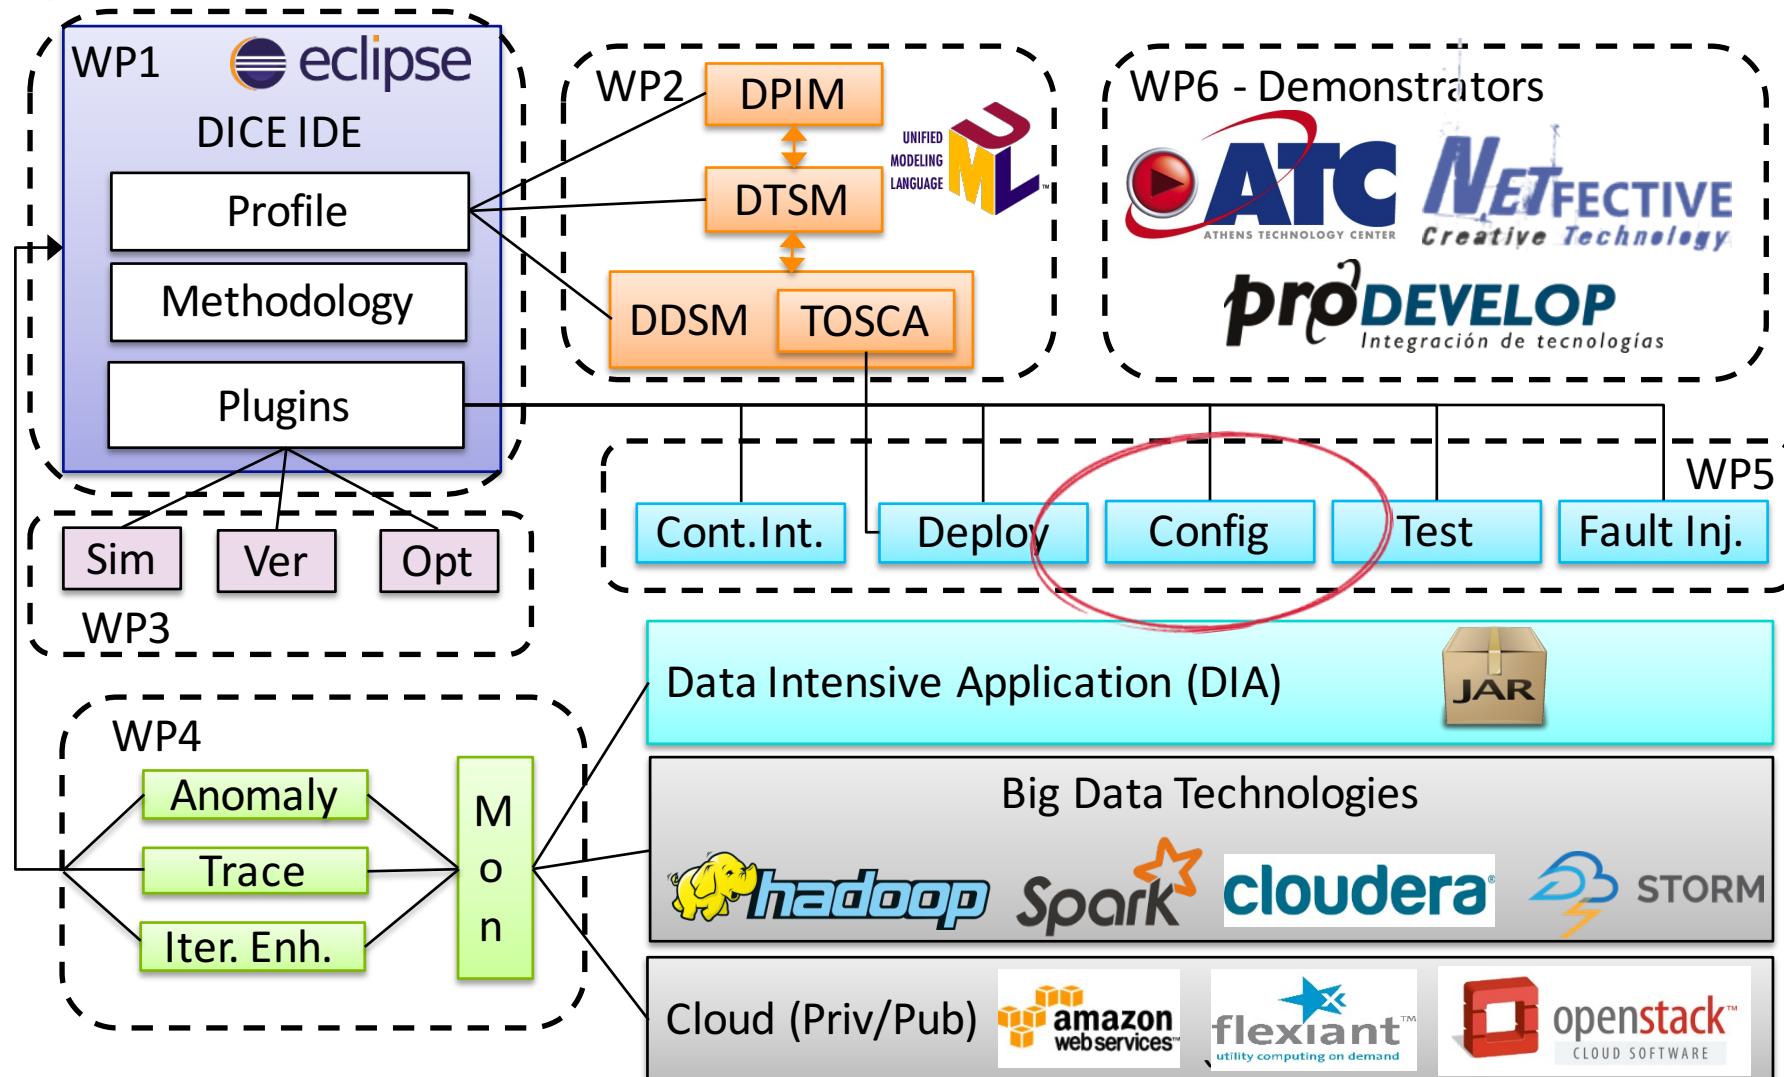
Runtime overhead



Key takeaways

- A principled way to leverage prior knowledge gained from searches over previous versions of the system.
- Multi-task GPs can be used in order to capture correlation between related tuning tasks.
- MTGPs are more accurate than STGP and other regression models.
- Application to SPSSs, Batch and NoSQL
- Lead to a better performance in practice

Acknowledgement: BO4CO and TL4CO are now integrated with other DevOps tools in the delivery pipeline for Big Data in H2020 DICE project



Tool Demo: <http://www.slideshare.net/pooyanjamshidi/configuration-optimization-tool>